

EducatedTraveler — The Guideline

The single source of truth for our line and how we develop.

This is the one document anyone reads before they make or build anything for EducatedTraveler — a caption, a clip, an Atlas note, an essay, a Circle issue, a feature, a script. The nine Parts that follow say *how* we do each thing; this front matter says *what we are and how to use the rest*. Start with The Line below, act from it, and only open the full Part when you need the detail behind a move.

Last reviewed: 2026-08-20

The Line (read this first)

One page. If you only read this, you can still make the right thing.

The ligne directrice. We map where real crafts are still alive, and we introduce people to the masters, the schools, and the peers already gathered there — so they can go learn something real, with their hands, with their people, at the source. The Atlas and the Circle are free — nobody pays to be on the map, and no money ever moves a line. One thing is paid: a Lab Week, built for the people of the Circle — and it never buys a place on the record. We are the bridge. Every line we publish has one job: pull a curious stranger to the next line, and finally toward the Circle.

The spine (say these). “A place, a person, your people.” · “Earned, not bought.” · “At the source, not the simulation.” · “Skills last, tans fade.” · The Atlas (the map) · The Circle (the community). Evoke freedom, adventure, immersion. Fight the screen; honour the source; people first. The customer is the hero; we are the guide.

THE ONE METRIC (next 90 days). Engaged Circle subscribers — sign-ups that actually open and click. Target ~150–300 quality sign-ups by day 90, 40%+ open. Not revenue, not followers, not pageviews. Every piece either feeds this number or it doesn’t ship.

THE FOCUS FILTER (checklist — any NO and it’s out, no matter the money on the table): - [] Embodied — hands on real material? - [] Immersive — a week or more, fully present? - [] Community-based — with peers, not solo? - [] Returns a real skill — you come home able to do a real thing? - [] Master test — does it fight the screen, or feed it?

Voice, in brief. Warm, plain, founder-led. No hype. The verb is INTRODUCE / CONNECT / BRIDGE — never sell, book, or enroll. Banned words: transformation, transformational, life-changing, vacation, holiday, luxury, easy. Public-surface bans: book, enroll, checkout, prices, “Quest”, XP/levels/badges, “choose your depth” as a pricing ladder. No decorative emoji (functional star/check/warn glyphs tolerated in UI only).

Design, in brief. Warm-Dark Editorial. Ink #0d0b09, surface #14110d, paper text #f3ede2 (never pure white). Accents: SEA #7fa8a5, EMBER #d28a52. Cores: wellness #94ad86 · adventure #6fa3a0 · creative #cf8f6e · culinary #c9a24a. Fraunces + Inter + IBM Plex Mono. Same warmth as Littoralicious, never its CSS or fonts. The cold blue/glass/gamified look is retired forever.

STOP-DOING. Don’t build more Atlas — 99 disciplines / 424 pages is plenty; distribution is the work now. Don’t re-surface selling on the public site. Don’t rebuild the open marketplace or instructor-availability portal.

Don't run a cold blitz or chase French CPF. Don't make "mastery vs the AI age" the public billboard (it's a quiet one-line foil). Don't pivot the Daily Drop format before day 90. Never cite an unverified market-size / TAM number. Don't reintroduce the cold look or copy Littoralicious CSS/fonts.

The 8-gate, in one line. No money and no curated immersion until all eight are true: (1) several hundred organic non-branded visitors/mo for 3+ months, (2) ~1,000–2,000 engaged Circle subs, (3) 40%+ open / 6–9% click sustained with replies and referrals, (4) regular unprompted "can you connect me?" inbound, (5) a few clips genuinely travelled, (6) at least one vetted partner willing to co-host on referral, (7) a soft refundable-deposit off-site immersion test that fills, (8) founder capacity to run it without starving Atlas + Circle.

Contents

1. Part 1 — The Editorial Line: how we make people *look*
2. Part 2 — Social Attraction: the angles
3. Part 3 — Instructors & Schools: the relationship program
4. Part 4 — The Community: a social home for skills
5. Part 5 — The Skill Passport: profile + credentials
6. Part 6 — IRL: connecting like-minded people
7. Part 7 — Development vectors: the different ways to grow
8. Part 8 — The Tram Rail: intelligent, honest automation
9. Part 9 — Sequencing & the focus filter

Appendix — One-page quick reference

Part 1 — The Editorial Line: how we make people *look*

This is the first page anyone reads before making anything for EducatedTraveler — a caption, a clip, an Atlas note, an essay, a Circle issue. Read it once. It governs every public surface so a stranger looks, trusts, and wants to go learn a real skill, hands-on, with their people, at the source.

One number tells you if it worked: **engaged Circle subscribers** — people who sign up and actually open and click. Not revenue, not followers, not pageviews. Every piece below either feeds that number or it shouldn't ship.

(a) The ligne directrice — read this before you make anything

We map where real crafts are still alive, and we introduce people to the masters, the schools, and the peers already gathered there — so they can go learn something real, with their hands, with their people, at the source. The Atlas and the Circle are free — nobody pays to be on the map, and no money ever moves a line. One thing is paid: a Lab Week, built for the people of the Circle — and it never buys a place on the record. We are the bridge. Every line we publish has one job: pull a curious stranger to the next line, and finally toward the Circle. We open on the payoff, keep one idea per piece, honour the source, and fight the screen. If a thing is embodied, immersive, with-people, and leaves you able to do a real thing — it belongs here. If it's a screen-feeder, a brochure, or hype, it doesn't, no matter how quick the win. We are the guide. The reader is the hero. "A place, a person, your people. Earned, not bought. At the source, not the simulation."

That paragraph is the whole job. Everything that follows is how to do it.

(b) Framing: make any skill read as embodied + with-people + at-the-source

The default way the world writes about learning is the trap. “A 5-day pottery course in Tuscany” reads as a product on a shelf — a *course* (a screen could give you that) or a *trip* (a tan that fades). Our framing has to make the reader feel the clay, the room, and the master’s hands before they feel a calendar or a cart.

The three moves that do it, every time:

1. **Embodied** — lead with the body and the raw material. Hands, heat, salt, grain, rope, stone. Never lead with a curriculum or a syllabus.
2. **With-people** — name the room. A circle, a cohort, the bar, the workshop, the locals already there. Never frame it as you-alone-improving.
3. **At-the-source** — anchor it to the one place on earth where this is most alive, and say *why* it’s alive there. Living capital, not a backdrop.

Before → After rewrites (use these as your ear):

Before (product / screen)	After (embodied · with-people · at-the-source)
“Learn to make sushi in a 7-day Tokyo course.”	“In a Tsukiji back-kitchen, a master watches your knife until the cut is right — beside others chasing the same edge. You don’t learn the rice; you learn the room.”
“Sailing holiday in Croatia — beginners welcome.”	“The Adriatic taught Europe to sail. You’ll be on the helm by Wednesday, the same channel the old crews ran, in a crew of strangers who’ll be your people by Friday.”
“Pottery retreat in Tuscany. Book your spot.”	“The wheel doesn’t lie. A week in a Tuscan workshop where the glaze recipes are older than the country — your hands, a teacher’s hands over yours, a table of people learning the same patience.”
“Wellness vacation in Kerala.”	“Kerala kept this practice alive for two thousand years. You go where the lineage still breathes, you learn it on your own breath, alongside people who came for the same reason.”
“Online breadmaking masterclass — lifetime access.”	“You can’t smell proof through a screen. In a Basque bakery at 4 a.m., the baker shows you the dough at the exact second it’s ready — a thing no video can teach.”

Notice the pattern: **a verb your body can feel + a place that earned it + a room you’re not alone in + a moment of competence at the end.** That last beat — “by Wednesday,” “until the cut is right” — is the payoff. It’s also our honest promise: *skills last, tans fade*.

Words that quietly break the frame (avoid even when not formally banned): *course, class, session, package, itinerary, getaway, programme*. They reattach the shelf. **Words that hold it:** *room, table, bench, helm, wheel, master, circle, source, by [day] you’ll be able to...*

And obey the bans without exception: never *transformation/transformational, life-changing, vacation, holiday, luxury, easy*. On any public surface, never *book, enroll, checkout, price, Quest, XP/level/badge*. We **introduce, connect, bridge** — we never sell, book, or enroll.

(c) The “worthy of their time & curiosity” bar — and how to open on the payoff

A stranger gives us one second. We earn the second one by being worth it. The bar is simple: **does this teach, surprise, or genuinely move someone toward going — or is it just noise asking to be liked?**

Earns the click (ship it): - A real fact they can carry to a dinner table tonight (the Gilda pintxo named after a Rita Hayworth film — true, cited, and it *means* something). - A hands moment so tactile it's almost rude not to watch (the Daily Drop's most tactile seconds, ending on the finished object). - A “wait, *that's* where that's from?” — the source revealed (the one town, the one lineage, the one reason it's alive there). - A person's honest “I came back able to...” — proof, not promise.

Noise (kill it): - Generic wanderlust (“paradise found”), aspiration with no skill inside it. - A listicle of “10 places to learn X” with no verified room behind it (this is the skillcation swamp — our biggest drift risk). - Anything that ends on a screen, a discount, or a vibe instead of a competence. - A fact we can't verify, dressed as fact (breaks the Honesty Lock — the moat).

Open on the payoff — always. The first sentence delivers the reward; it does not warm up to it. Our own live essay is the model: “*Here's something you can carry to your next dinner table.*” Then the Gilda. Then the lesson. Never bury the gift under a throat-clear.

- **Bad open:** “San Sebastián is a beautiful coastal city with a rich culinary tradition...”
- **Right open:** “The first pintxo ever named was named after a Rita Hayworth film. Here's why that matters for anyone learning anything well.”

If you can't write a first line that gives something away, the piece isn't ready.

(d) The four content pillars — what each is FOR, and how it feeds the Circle

Four pillars, one job each. Each one ends by handing the reader to the Circle. (The **Daily Drop** is the locked top-of-funnel engine that *carries* these — do not pivot it pre-day-90.)

1. The Atlas — through the meaning lens. *For:* proving the source is real. The Atlas isn't a directory; it's an argument that this craft is *alive here, with these people, and you can go*. We write each entry through the meaning lens — not “what's here” but “why this place earned this craft, and where your cohort gathers.” *Feeds the Circle:* every Atlas page is a destination for the Daily Drop's one URL; the meaning hook (“Living Capital,” “Strongest Scenes”) is what makes a viewer want more than one place — and the only honest way to get “more” is to join the Circle. *Stop-doing reminder:* the Atlas is **done** at 99 disciplines / 424 pages. Distribution is progress, not page count. We deepen meaning and verify; we don't add bulk.

2. Person-led “I came back able to...” stories. *For:* turning promise into proof. One person, one place, one real skill they now have. First-person, plain, specific — what their hands could do on day one vs. day seven, who was in the room with them. *Feeds the Circle:* this is the highest-trust asset we have and the most direct driver of the one metric. It answers “is this real?” with a face. End every story on the Circle, not a hard ask — “*This is the kind of room we map. The Circle is where the next one gets found.*” *Honesty caveat:* never fabricate or composite a person. No verified story yet → don't manufacture one; lean on Atlas + Journal until a real one (or a travelled clip) exists. This is a now-gather, not a now-invent.

3. The Journal — science-and-stories essays. *For:* depth and authority. The reason a craft is alive where it's alive — history, science, lineage, the quiet lesson underneath. Warm, cited, JSON-LD, 4-minute read. San

Sebastián is the template: a fact you can carry, then the meaning, then *the same lesson about learning anything well*. *Feeds the Circle*: the Journal earns the *trust* half of “look, trust, want to go.” It’s our SEO and our credibility, and it converts the slow reader — the one who needs to know we’re serious — into a subscriber. Every essay closes by pointing to the Atlas page it lives near and the Circle. *Bar*: if there’s no verifiable fact and no real lesson, it’s a blog post, not a Journal essay. Cut it.

4. Put-the-screen-down pieces. *For*: the quiet foil. Short, sharp, honest reminders that some things can’t be learned through glass — “*you can’t learn a roll-tack from a video*.” This is our differentiation from MasterClass/YouTube, made by feel, not by argument. *Feeds the Circle*: it reframes the reader’s restlessness with screens into a reason to go *do* — and the Circle is the door to doing. *Stop-doing reminder*: “embodied mastery vs. the AI age” is a **quiet one-line foil, never the billboard**. The market has buzzword fatigue. Drop it in; never build a piece around it as a manifesto.

The shared rule across all four: every pillar ends pointing somewhere — Atlas page, Circle, or both. A piece that doesn’t hand the reader forward isn’t finished. That hand-forward is the only conversion mechanism we have, and it converts into the one metric.

(e) The curiosity engine — open a loop honestly, never clickbait

Curiosity is the engine; honesty is the brake that keeps it ours. An honest loop **promises something true and then delivers it**. Clickbait promises and withholds, or promises what isn’t there. We open loops we can actually close.

How to open a loop honestly: - **Name a real gap the reader feels**: “You can watch a thousand pottery videos and still not know when the clay is ready. There’s a reason.” → then *give* the reason or the place that teaches it. - **State a true surprise, then pay it off**: “The world’s greatest eating city is built on standing up. Here’s what that taught me about learning.” (We pay it off — that’s our live essay.) - **Use the Daily Drop’s wordless tension honestly**: the clip is the loop — hands working raw material — and it closes on the finished object. The competence is the payoff, on screen. No bait; the answer is in the video.

The honest-loop test (must pass all three): 1. Is the promise *true*? (Honesty Lock — every fact hand-verified.) 2. Do we *deliver* the payoff in this same piece, or hand it to a real next step (Atlas/Circle)? 3. Would the reader feel *rewarded*, not tricked, after reading?

Banned curiosity moves: “You won’t believe...”, manufactured outrage, fake scarcity, “the one weird trick,” withholding a payoff that doesn’t exist, any TAM/market-size number (the \$75B–\$432B figures were refuted — never cite them). One betrayed loop costs us the trust that is the business.

A note on numbers, everywhere we make things. Any performance figure you find in 2026 marketing or platform write-ups (engagement multipliers, reach percentages, “best mix” ratios) is from third-party analyses — directional, not gospel. Verify it in the Studio Metrics tab against our own opens and clicks before we ever bet a format on it. Our moat is honesty; never restate someone’s blog stat as a settled fact in our copy.

(f) The focus filter — as an editorial checklist

Run this on every immersion, partner, article, post, clip, Atlas note, and feature *before* it ships. **Any NO = it's out**, regardless of how quick the money or the reach.

- ☐ **Embodied?** — hands and real material are in it (not a concept, not a screen).
- ☐ **Immersive?** — it points to being fully present (a week+, a real room), not an afternoon or a feed.
- ☐ **Community-based?** — there are *peers* in the frame, not a solo improver.
- ☐ **Real-sense transformational?** — the reader/learner comes back able to *do a real thing*. (Internal label only, in our canon Focus Filter vocabulary. We feel it; we never use the word — and this checkbox never gets copy-pasted onto a public surface.)
- ☐ **Master test: does it fight the screen or feed it?** — if it just feeds the scroll, kill it.

Then the editorial overlay:

- ☐ **Opens on the payoff?** — first line gives something away.
- ☐ **One idea?** — one piece, one point; everything else cut.
- ☐ **Re-hooks and closes on one thing?** — pulls the reader through and lands on a single takeaway.
- ☐ **Verified?** — every fact hand-checked against the source; anything shaky marked “provisional, verifying,” never stated as fact.
- ☐ **Bridge voice?** — introduces/connects/bridges; never sells/books/enrolls; no banned word; customer is the hero, founder is the guide.
- ☐ **Hands forward?** — ends pointing to the Atlas, the Circle, or both.
- ☐ **Moves the one metric?** — a plausible path from this piece to an engaged Circle subscriber. If not, it's noise.

If a piece passes the focus filter but fails “hands forward” or “moves the one metric,” it's a nice thing that doesn't serve the mission — fix the ending or don't ship it.

Now-moves vs. later-moves (so this page stays honest about phase)

- **Now:** Atlas-through-meaning, Journal essays, put-the-screen-down pieces, the Daily Drop, and gathering the first real “I came back able to...” stories as they happen — all driving the **one metric** via the Circle. No selling, ever, on the public site.
- **Later (Phase-2 gated, all 8 gates true first):** anything that frames a *specific bookable* immersion, partner co-host, or money. Until then we frame the *craft and the source*, never a transaction — and we keep the dormant selling machinery dormant.

The whole page in one line: make a stranger feel the hands, trust the source, want their people — and walk them to the Circle. Skills last, tans fade.

Part 2 — Social Attraction: the angles

The job of this system

The Daily Drop earns the eyeball. This system earns the *follow, the save, the share, and the sign-up* — and routes every one of them to a single destination: an Atlas page or the Circle. Nothing here pivots the Daily Drop.

Everything here LAYERS around it.

One metric governs the whole thing: **engaged Circle subscribers** (target ~150–300 quality sign-ups by day 90, 40%+ open). A post that goes viral but feeds zero sign-ups is a failure here. A quiet carousel that puts twelve real people on the list is a win. Every angle below is judged on that, and on one master test: **does it fight the screen or feed it?**

*Read the numbers below as directional, not gospel. Every figure in this Part comes from third-party 2026 marketing/SEO analyses — useful for shaping the format, but not platform-published laws. Treat them as hypotheses to validate in the **Studio Metrics** tab against our own per-stream data before betting the format on any of them. Our moat is honesty; we don't repeat other people's claims as if we measured them ourselves.*

The 2026 reality we're building on, per 2026 third-party analyses (directional — verify in Studio Metrics): - **DM shares look like the kingmaker.** Marketing analyses suggest shares carry outsized weight for Reels distribution — one analysis puts it at roughly 3–5x a like — and rank watch time, likes-per-reach, and DM shares as the top signals (Later, 2026; Mirra, 2026). Treat the multiplier as a claim, not a constant. Either way, our angles are engineered to be *sent to a friend*, not just liked. - **Niche specificity beats reach.** The tighter the niche, the better the algorithm targets genuinely interested viewers — these guides argue sustainable growth comes from serving a specific community, not everyone (WordStream, 2026). ET is already maximally niche: real craft, at the source, with peers. - **Raw beats polished.** Founder-led, behind-the-scenes, candid is reported to out-perform over-produced ads; viewers prefer realness over perfection (WordStream, 2026). This is our home turf — Arnaud is real, the places are real. - **Carousel = conversion, Reels = discovery.** One 2026 guide recommends a mix of roughly 60–70% Reels, 20–30% carousels, <10% static, and reports carousel engagement up around 30% — calling carousels “your conversion engine” (Second Brain, 2026). Take the split as one guide's recommendation, not a rule; tune it in Studio Metrics. - **Link-in-bio friction is a documented conversion cost**, and end-card CTAs are reported to lose people before the final frame (Influencers-Time, 2026). So the funnel below removes friction, not adds steps. - **Pinterest = compounding evergreen traffic**, reportedly strongest for visual niches (travel/craft/food); these analyses claim a well-built pin can drive clicks for many months and that the large majority of top searches are unbranded (Pin Nerds, 2026; Improvado, 2026). If even directionally true, that is the Atlas's natural air supply.

(a) The angle bank — 16 distinct streams

The Daily Drop is Stream o (locked, untouched). These sixteen are the *layers around it*. Each is built to fight the screen, route to one destination, and earn a save or a share. The existing seven hooks (The Source / Living Capital / Strongest-Community / Before→Competence / Can't-do-from-a-screen / The People Here / The Number) live *inside* the Daily Drop captions; the streams below are whole content types beyond it.

Format legend: **R** = Reel/Short, **C** = carousel, **P** = Pinterest pin, **L** = long-form, **J** = Journal/article, **T** = text post (Threads/X/IG note). Core legend: Wellness #94ad86 / Adventure #6fa3ao / Creative #cf8f6e / Culinary #c9a24a.

#	Stream	Hook pattern	Why it fights the screen	Best format · core · platform	Routes to
1	The Source Map	“This is the one place on earth where [craft] is still alive.” A slow pan of a real workshop/coast/kitchen, then the place-name.	Anchors a skill to a <i>real geography</i> you’d have to travel to — the opposite of “anywhere on a screen.”	C + P · all cores · IG, Pinterest	Atlas page for that craft
2	Master’s Hands	Tight on the master’s hands doing the one move a beginner can’t fake. No face needed.	Pure embodied craft; the hands ARE the argument.	R · Creative/Culinary · TikTok, Reels, Shorts	Atlas school card
3	The Circle Around the Fire	A wide shot of the <i>cohort</i> — strangers who became people-chasing-the-same-fire. “VAWAA pairs you alone with one artist; here you’re in a circle.”	Shows the peer-community no screen and no solo workshop can give.	R + C · all cores · Reels, TikTok	Circle sign-up
4	Before → Competence	Day 1 wobble vs. Day 6 the-thing-works. Same person, same hands.	Proof that a week+ at the source actually returns a real skill — “skills last, tans fade.”	R · Sportive/Adventure · Shorts, Reels	Atlas page
5	The Honest Atlas	“We checked this against the school’s own site. Verified.” Show the actual source line.	Honesty as spectacle — the moat made visible; counter to the listicle swamp.	C + T · all cores · IG, Threads	Atlas page
6	The Number	One verified figure on a clean card: “4 masters left who still do this by hand.”	A single fact is inherently save-able and send-able.	C + P · Creative/Culinary · Pinterest, IG	Atlas page
7	What Do You Want to Become?	The five modes as a question, not a catalog: Adventurous · Creative · Healing · Sportive · Crafting.	Reframes “what class” → “who you become” — identity, not product. (Homepage tiles read as identity, never as a depth/pricing ladder.)	C · all cores · IG, Pinterest	Homepage / Circle
8	Field Notes from the Founder	Arnaud to camera, at sea or in a real kitchen: one honest reason a place made the Atlas.	Candid founder content is reported to out-perform polished	R + T · all cores · TikTok, Threads, Shorts	Circle sign-up

#	Stream	Hook pattern	Why it fights the screen	Best format · core · platform	Routes to
			— and it's <i>true</i> , because he's there.		
9	The Quiet Foil	One line, no billboard: “You can’t learn a roll-tack from a video.” Over real footage of the real move.	Names the screen, then shows the source — the foil stays a whisper (canon: not the public billboard).	R · Adventure/Sportive · TikTok, Reels	Atlas page
10	When to Go	“If you want [craft] alive and the season right, you go in [month].” A calendar truth.	Practical, evergreen, deeply save-able — and only exists because someone went.	C + P · all cores · Pinterest, IG	Atlas When-to-go
11	Two Crafts, One Coast	“In this one region you could learn [X] and [Y] within a day’s drive.”	Surprising adjacency = a “send this to your travel friend” moment.	C + P · Adventure/Creative · Pinterest, IG	Atlas Region lens
12	The Strongest Scene	“If you want a <i>cohort</i> , not a private lesson, this is the strongest scene for [craft] right now.”	Community-RANK made tangible — best chance to find your people.	C + R · all cores · IG, Reels	Atlas Strongest-Scenes lens
13	Ask the Atlas	Reply to a real DM/comment (“where do I actually learn X?”) with the verified answer on screen.	Turns the publication into a <i>responsive human</i> — replies are the engaged-Circle behaviour we want.	R + T · all cores · TikTok, Threads	Atlas page / Circle
14	The Journal Pillar	A real essay (science-and-stories), teased as a 3-card carousel: claim → evidence → “the place this is true.”	Publication-grade depth no skillcation listicle can match; SEO + authority.	J + C · all cores · web, IG, Pinterest	Journal → Circle
15	Material Truth	Extreme close-up of the <i>raw material</i> — wet clay, line tension, dough at the turn — then what it becomes.	Tactility you can almost feel; the “embodied” filter as pure sensation.	R · Creative/Culinary · TikTok, Reels	Atlas page
16	Who’s Already There	Anonymous-but-real: “Three people on the list are eyeing [place] this autumn.” (only when true)	Real, gentle social proof of momentum — never fabricated, never	T + C · all cores · Threads, IG	Circle sign-up

#	Stream	Hook pattern	Why it fights the screen	Best format · core · platform	Routes to
			countdown-timer scarcity.		

Sample ET-voice captions (warm, plain, founder-led; customer is the hero):

- **Stream 1 (The Source Map):** “There are coasts where boatbuilding is a memory and one where it’s still a morning’s work. This is the second kind. The masters here still bend oak by hand, and a few of them teach. A place, a person, your people. The Atlas keeps the address. [atlas URL]”
- **Stream 3 (Circle Around the Fire):** “None of these people knew each other on Monday. By Friday they were finishing each other’s stitches. That’s the part a private lesson can’t give you and a screen never will. Earned, not bought. [Circle URL]”
- **Stream 5 (The Honest Atlas):** “Before this went in the Atlas, we read the school’s own page and matched every claim. If we couldn’t verify it, we marked it provisional and kept digging. That’s the whole job. [atlas URL]”
- **Stream 8 (Field Notes):** “I’m writing this from the galley, somewhere off the coast. I added a place to the Atlas today because a chef there still ferments the old way and welcomes a small group each spring. I went looking so you wouldn’t have to guess. [Circle URL]”
- **Stream 13 (Ask the Atlas):** “Someone asked where you actually learn to throw on a kick-wheel — not watch, *learn*. Here’s the verified answer, and the season to go. Ask, and the Atlas answers. [atlas URL]”

Caveat: Streams 3 and 16 require *real* cohort footage and *real* list movement. Until you have shot a genuine cohort, Stream 3 runs on archival/partner footage clearly labelled — never stock posing as your community (see the never-post list).

(b) Per-platform play, keyed to 2026

The Daily Drop already goes same-clip/same-day to TikTok + Reels + Shorts on @educatedtraveler.app. The streams above slot into each platform by its 2026 strength. (Per 2026 third-party analyses, directional — verify in Studio Metrics.)

Instagram (Reels + carousels) — the conversion home. One 2026 guide recommends a mix of roughly 60–70% Reels, 20–30% carousels, <10% static ([Second Brain, 2026](#)) — start there, then let Studio Metrics move it. Reels (Daily Drop + Streams 2, 4, 9, 15) bring new eyes via Explore; carousels (Streams 1, 5, 6, 7, 10, 11) do the converting — that same guide reports carousels earn around 10–12% more engagement and serve as the natural save/SEO surface. Engineer every Reel for a **DM share** (the signal those analyses weight most): end on a moment so satisfying the instinct is “send this to [name].” Use 3–5 niche hashtags only — more is reported to reduce reach ([Later, 2026](#)). Caption SEO matters: write the craft + place in plain words early. Now-move.

TikTok — discovery and founder-voice. Best for Master’s Hands (2), Material Truth (15), Quiet Foil (9), Field Notes (8), Ask the Atlas (13). Candid founder content is reported to out-perform polished here ([WordStream, 2026](#)). Front-load the most tactile second; native sound, not licensed music, keeps it honest and ambient (canon). Caveat: link-in-bio friction is documented as real on TikTok ([Crescitaly, 2026](#)) — so TikTok’s job is *reach* + *profile visit*, and the bio does the routing. Now-move.

YouTube — Shorts now, long-form as a phase-gated lever. Shorts take the same Daily Drop + Before→Competence (4) + Strongest Scene (12). **Long-form (Stream 14-adjacent) is a later-move:** a 6–

10 min “Is this craft still alive here?” mini-doc is exactly ET’s publication-grade lane and YouTube’s evergreen search is widely reported as unmatched — but it’s founder-capacity-heavy. Gate it: don’t start long-form until the Daily Drop cadence is stable past day 90. Hub the link-in-bio to one Circle landing (Linkero, 2026). Shorts = now; long-form = post-90.

Pinterest — the compounding Atlas air supply. This is the single highest-leverage *non-video* now-move. These analyses claim a well-built pin keeps driving clicks for many months and that the large majority of top searches are unbranded (Pin Nerds, 2026) — meaning, if directionally true, people searching “where to learn weaving in person” find the Atlas without ever knowing the brand first. Pin Streams 1, 6, 10, 11, 14 with keyword-rich titles/descriptions, one board per core + one per region, each pin linking *directly to its Atlas page*. Seasonal lead time is reported as real: to own “autumn pottery retreat” you likely need to be indexed by summer (Improvado, 2026). Now-move, high priority — and the cleanest place to test those evergreen claims against our own Studio Metrics.

Where a publication-grade brand wins (the moat surface). The Journal + the Atlas itself are the defensible ground no skillcation listicle or funded incumbent can copy on voice and honesty. Threads/X carry text-first streams (5, 8, 13, 16) — low-cost, founder-real, conversational, and the place “Ask the Atlas” replies live. This is where ET *sounds like a publication, not a feed*. Now-move, low effort.

(c) Authenticity rules — what keeps it worthy

These are non-negotiable. They are the moat. Break one and the honesty edge over the listicle swamp is gone.

1. **Show the real place / real person.** Every clip is shot where the craft is actually alive. If Arnaud hasn’t been and there’s no verified partner footage, it doesn’t post.
2. **No stock-footage lies.** Stock or archival is allowed ONLY when clearly labelled and never implied to be ET’s own cohort or master. Posing models = banned.
3. **Founder-real.** Field Notes are Arnaud, actually at sea or in a real kitchen, in his own voice — not a script read by a stand-in.
4. **No fake scarcity.** No countdown timers, no “only 2 spots” theatre, no manufactured urgency. Stream 16 runs only on *true* list movement and stays gentle.
5. **No clickbait.** The hook must be delivered by the post. If the caption promises “the last place on earth,” the post proves it or the caption changes.
6. **Verified or marked.** Any Atlas fact on screen is hand-verified against the school’s source; unverifiable claims are marked “provisional, verifying” — on social too, not just on the site. The same discipline applies to the marketing stats in this Part: when we cite a number publicly, we cite it as a claim with its source, never as our own measured fact.
7. **No price talk, ever, on public surface.** No prices, no “book,” no “enroll,” no checkout language. The verbs are introduce / connect / bridge.
8. **No decorative emoji.** Functional glyphs (star/check/warn) only, UI-side. Reads as AI otherwise.

The never-post list: - Stock footage dressed as our cohort or our master. - Any countdown timer, “spots left,” or fabricated urgency. - A skill taught from a screen presented as the real thing (we *contrast* the screen; we never sell one). - Prices, payment, “book/enroll/checkout,” or any dormant selling machinery (Quest, XP, badges, “choose your depth” ladder). - An unverified market-size / TAM number (the \$75B–\$432B figures were refuted — never cite). - The “embodied mastery vs the AI age” line as a billboard (quiet one-line foil only). - Banned words: transformation/transformational, life-changing, vacation, holiday, luxury, easy. - A claim we

couldn't verify, stated as fact (including a borrowed marketing stat dressed up as our own finding). - Anything that fails the Focus Filter: not embodied, not immersive (week+), not community-based, no real returned skill.

(d) Curiosity / saveable / shareable mechanics (no dark patterns)

The goal is to earn saves and DM shares *honestly* — and DM shares are the distribution signal those 2026 analyses weight most heavily (directional; confirm in Studio Metrics whether shares actually move our reach) (Mirra, 2026).

Saveable by design (carousels + pins — reference value people keep): - **The single verified Number** card (Stream 6) — one fact, clean, keep-worthy. - **When-to-go calendars** (Stream 10) — practical reference, the textbook save. - **Two-Crafts-One-Coast** maps (Stream 11) — trip-planning utility. - **The five modes** as a “which are you?” card (Stream 7) — identity reference. - Frameworks/checklists are reported to perform well (TrueFuture, 2026) — but ours are real (“the 6 questions that prove a craft is still alive here”), never filler.

Shareable by design (the “send to a friend” instinct): - End every Reel on the *moment of competence* — the finished object, the move that lands. That satisfaction is what gets it sent. - “Two-crafts-one-coast” and “strongest scene” naturally tag a specific friend (“this is your trip”). - Honest surprise (a real Number, a real verification) is sent because it's *true and rare*, not because of a bait headline.

Honest curiosity (open-loop without lying): - Open on the payoff, then deepen — “This is the last coast where they still bend oak by hand” *and the post proves it*. The curiosity is real because the answer is real. - “Ask the Atlas” turns a genuine question into the post — curiosity the audience supplied.

Explicitly banned mechanics: engagement-bait (“comment YES”), follow-to-see-the-rest, false cliffhangers the post never resolves, rage-bait, manufactured controversy, and any “link in bio” tease for content that doesn't exist.

(e) The funnel and the one-CTA discipline

post → profile → Atlas page or Circle → engaged subscriber.

Reel/Short (reach) → profile visit → link-in-bio hub → Atlas page → Circle sign-up (the metric)
Carousel/Pin (intent) → direct link → Atlas page / Journal → Circle sign-up
Field Note / Ask-the-Atlas (trust) → profile → Circle landing → engaged subscriber

These guides converge on one rule: one primary CTA above the fold (Influencers-Time, 2026). For ET:

One CTA per post. Pick the destination by stream, never stack two. - Discovery streams (Reels/Shorts: 2, 4, 9, 15) → CTA is *profile* → *Atlas page*. Don't ask for the email mid-discovery; earn the visit first. - Intent streams (carousels/pins: 1, 5, 6, 10, 11) → CTA links *straight to the Atlas page* (skip the bio friction where the platform allows the direct link — friction is the documented conversion cost, Crescitaly, 2026). - Trust streams (Field Notes, Circle, Strongest Scene, Who's-already-there: 3, 8, 12, 16) → CTA is *join the Circle*.

Build target — the link-in-bio hub. This is on the “ABSENT TODAY” list and is a now-move: a Linktree-style ET-branded landing (Warm-Dark Editorial, not Littoralicious CSS) with exactly two doors above the fold — “Open the Atlas” and “Join the Circle” — plus the latest Journal piece. One hub, two doors, no clutter

(Linkero, 2026). The Atlas page itself must carry the Circle email capture (#circle → launch_waitlist) so an Atlas visit can become a sign-up in one step — the Atlas is the bridge, the Circle is the destination.

CTA language discipline (verbs are introduce / connect / bridge): - Yes: “The Atlas keeps the address.” · “Join the Circle.” · “Ask, and the Atlas answers.” · “I went looking so you wouldn’t have to guess.” - Never: “book now,” “enroll,” “checkout,” “limited spots,” “sign up before it’s gone,” any price.

The single funnel KPI is not reach — it’s Circle sign-ups that open and click. A stream that posts beautifully and produces zero sign-ups gets cut or re-routed in the Studio Metrics tab. Track per-stream: profile-visit → Atlas-page → Circle-signup, and weight the angle bank toward whatever actually feeds the 40%-open list. This is also where the borrowed 2026 stats above get tested against reality — keep what our own numbers confirm, drop what they don’t.

Now vs. later (honour the stop-doing list)

Now-moves: all 16 streams as content (built on the live Daily Drop + Atlas + Journal + Circle); Pinterest as the evergreen Atlas air supply (highest-leverage non-video); the two-door link-in-bio hub; Threads/X text streams; per-stream CTA routing in the Studio.

Phase-gated later-moves: YouTube long-form mini-docs (post-90, capacity-gated); any partner/co-host content (only after the 8-gate); the skills-only social feed, public Linktree-profiles-for-members, and verifiable credentials (Atlas-adjacent builds, *not* social-attraction, and not before the Circle metric is hit).

Do not, ever: pivot the Daily Drop before day 90; re-surface prices/booking/Quest/XP on any public post; build more Atlas pages to “do social” (distribution is the progress); cite a TAM number; make the AI-foil the billboard; present a borrowed marketing stat as a settled fact; copy Littoralicious fonts/CSS or the retired cold-blue look.

Every angle here exists for one reason: to put a real person on the list who opens, clicks, and one day asks, “can you connect me?” That ask is the whole point. Skills last, tans fade.

Part 3 — Instructors & Schools: the relationship program

The Atlas already names 835 verified schools. This program is how a name becomes a relationship — and how a relationship, over many months, becomes a clip that travels, a school that reposts us, and (only at Phase 2) one or two partners ready to co-host. The whole thing runs on one promise we never break: **we introduce, and the vouch is never for sale.** A school that gets a beautiful feature and a wave of warm strangers from us learns to trust us long before money is ever on the table. That trust is the asset. The metric it feeds is the only one that matters now — **engaged Circle subscribers** — because every school we feature is a school that shares our clip, and every share puts our “what we are” line and our one Atlas URL in front of exactly the people most likely to open and click.

One discipline at a time. Everything below passes the Focus Filter first: if a school isn’t embodied, immersive, community-based, and real, it isn’t on the Atlas, so it never enters this program.

Read this before you trust any number below. The platform stats, feature dates, gear prices, and standards-status dates in this section are drawn from third-party 2026 analyses, vendor pages, and marketing write-ups — directional, not gospel. Treat them as starting points to verify in the Studio Metrics tab (for reach claims) or at build time (for prices, feature dates, and standards status) before you bet a format or a budget on them. Our moat is honesty; we don't launder a blog claim into a settled fact.

(a) The outreach ladder — seven rungs, what we offer, what we never do

Think of it as a slow staircase, not a funnel. Most schools will sit happily on rung 3 or 4 forever, and that's a win. Nobody is pushed up. **At every rung the offer is the same three things: distribution (real eyes from people who want this craft), a beautiful feature (we make them look like the masters they are), and trust (we never ask them to pay, and we say so out loud).** We never raise the immersion business until the Phase-2 gate is met — and even then, only with the handful who've climbed on their own.

Rung	Name	What ET does	What ET offers	What we NEVER do
0	Cold find	Hand-verify the school against its official source; add a confidence-tagged Atlas card	(nothing yet — the card links them for free)	Fabricate, rank by who pays, or list a school we haven't verified
1	First contact	Send the “you're on the Atlas” note — a gift, not an ask	A live, linked feature already published	Lead with a request; mention money, cohorts, or fees
2	Permission to feature	Ask to share 1–3 of <i>their</i> existing clips, fully credited	Distribution to a craft-hungry audience; we do the editing	Promise placement, traffic, or ranking
3	Footage licensing	Get a written, scoped license for the clips	A finished film in the ET look they can reuse	Use a clip before written yes; alter terms silently
4	Repost partner	They reshare our cut of their work to their own followers	Two-way audience flow; a recurring beautiful frame	Demand reciprocity; make resharing a condition
5	Co-created video	We (or they, on our brief) shoot something new together	A founder-grade short film, free, theirs to keep	Treat it as a sales deliverable or paid ad
6	Phase-2 co-host (gated)	Invite to co-host one curated immersion, off-site, on referral	A warm community already chasing their craft	Open this until all 8 Phase-2 gates are green

Voice rule for the whole ladder: the school is the hero, Arnaud is the guide. Warm, plain, founder-led. The verb is always *introduce* / *feature* / *connect* — never *sell* / *book* / *enroll*. No hype words, ever.

Rung 1 — First contact (the gift email)

This is a *new* note, distinct from the licensing ask. It arrives with nothing to give back. Send from arnaudcallier@pm.me. Track every send in the Studio Outreach tab. Note: rung-1 outreach (and rungs 2–5)

NEVER references co-hosting, cohorts, or immersions — that line lives only at rung 6, and only after the gate. Keep it sealed off here.

Subject: {SCHOOL} is on the EducatedTraveler Atlas

Dear {NAME},

I'm Arnaud Callier — a French chef, fifteen years working on the water, and the founder of EducatedTraveler. We keep a hand-verified map of the places where a real craft is still truly alive, ranked by the strength of the community that gathers there. No bookings, no commissions, no prices. We introduce people to the right school and then get out of the way.

I wanted you to know {SCHOOL} is featured on our {DISCIPLINE} page for {PLACE}, with a link straight to you: {ATLAS_URL}

Nothing's needed from you. I just believe people looking for {DISCIPLINE} should find you, and now they can.

Warmly, Arnaud EducatedTraveler — a place, a person, your people.

That's the whole email. No ask. It earns the reply that opens rung 2.

Rung 2–3 — Permission + footage licensing

The existing **OUTREACH.md** email is the right tool here and stays the workhorse — it sources the daily clip *and* quietly opens the partner door over time. Keep using it, with one tightening: split the soft “I'd also welcome a call” line into its own follow-up so the first email is *purely* a feature, never a two-job email that reads as a setup. The call line moves to rung 5/6, after a clip has actually travelled. Even though OUTREACH.md doubles as partner recruitment in our own heads, its rung 1–5 wording must read as a feature and nothing more.

Rung 4 — Repost partner (the ask that isn't an ask)

After we publish their cut, we send the finished film *first*, then make resharing effortless and optional:

Subject: Your {DISCIPLINE} film is live

{NAME} — here's the short we made from your footage: {LINK}. It's yours — post it anywhere, anytime, no need to credit us.

If it feels right to share with your own people, here's the ready version for Reels/TikTok/Shorts: {ASSETS}. We've credited and tagged you in every caption. Either way, thank you for letting us show your work.

Rung 5 — Co-created video

Only for schools that have replied warmly and reshared at least once. This is where Arnaud's media moat shows: we offer to make something *new* together — either he shoots it on a port call, or the school shoots to our remote kit and brief (below). It's a gift, not a deliverable. The “short call” line from old OUTREACH.md belongs here, reframed:

I'd love to make one proper short film with you — your hands, your place, the craft as it really is. I'll either come film it or send you a one-page kit so it's effortless. No cost, no catch. It's the kind of thing we make for the schools we admire most.

Rung 6 — Phase-2 co-host (DO NOT OPEN EARLY)

This is the *only* rung where money is named, and it's a **later-move, hard-gated by all 8 Phase-2 conditions**. It must never appear on the public site — this copy is internal-only. Until the gate turns green, the shortlist just quietly fills (see section d). When it opens, the frame is still connection: "We've got a circle of people chasing exactly your craft and asking us to introduce them. Would you be open to hosting a small group, once, on a referral arrangement we'd handle off-site?" Never a contract-first pitch.

Now-moves: rungs 0–5. **Later-move (gated):** rung 6. Do not let the existence of rung 6 leak into rungs 1–5; that's how trust dies.

(b) The content-sharing engine — repost their clips, our frame, clean rights

The daily clip is the engine that feeds the one metric. Schools are our richest source of footage *and* our biggest distribution multiplier when they reshare. Here's how to run it without ever crossing a rights line.

Rights & permission hygiene (the 2026 norms, applied)

A repost is not consent and a tag is not a transfer of rights — copyright vests automatically in whoever shot the clip, and using it commercially without an explicit, written, *scoped* grant exposes you to real liability ([Showcase](#), [Later](#)). So ET's rule is simple and non-negotiable: **no clip is touched before a written yes that names the boundaries**. A defensible license names platforms, duration, whether edits are allowed, and geography ([Digital Applied](#)). Our reply-thread grant does exactly that, in plain founder English:

Footage permission — plain terms (paste into the email reply): *By replying "yes," you give EducatedTraveler a non-exclusive license to edit and post the clip(s) you send — short films on our social accounts (TikTok, Instagram, YouTube, our site) and our newsletter, worldwide, for as long as we both want. You keep full copyright; you can post the clips anywhere yourself. We always credit and tag you. No paid ads without asking you first. Any music we add is licensed by us. If anyone recognisable is in frame, please confirm they're fine being shown.*

That single paragraph covers scope, duration, modification rights, geography, non-exclusivity, paid-ad carve-out, music clearance, and likeness clearance — the full 2026 checklist ([Showcase](#), [The Social Media Law Firm](#)). Keep every "yes" thread; the Studio Outreach tracker logs *date, school, scope, link to thread*. Turn the messy DM into a repeatable record — capture request, store consent, publish with attribution ([Automateed](#)).

(One caveat on all of the above: these are guidance pages and marketing/legal write-ups, not legal authority. The plain-terms grant is a sensible default, not a substitute for a real review — before we lean on it at any scale, have a lawyer look at it. General guidance, not legal advice.)

Two modes, two rights levels: - Reposting their public clip as-is (Instagram's native repost — reported live since around Aug 2025; verify the current feature/date before you rely on it) — light use, still ask first, credit

visibly in the caption and tag them in-frame; burying credit in hashtags doesn’t count (Later). - **Re-editing their footage into an ET film** — this is the licensed path above; needs the written grant because we’re modifying and distributing off their platform.

Attribution rules (always, every time)

- 1. **In the caption, first line of credit:** “Footage: @school, {place}.” Visible, not in hashtags.
- 2. **Tag the school** in the post and in-frame where the platform allows.
- 3. **Atlas link** to their page is already live — the caption’s one URL points there, sending warm traffic both to them and to the Circle sign-up.
- 4. **Their cut is theirs:** when we send the finished film, they may repost with no obligation to credit us back. Generosity compounds trust.
- 5. **If a school ever asks to be removed** — clip down within 24h, no questions. Honesty Lock.

The consistent ET frame (so a school’s clip still reads as “EducatedTraveler”)

Every shared clip — ours or theirs — is wrapped in the same locked Daily Drop format so the feed has one face: 30–60s vertical, **wordless**, hands + raw material, the most tactile seconds, **end on the finished object / moment of competence**, one place-name caption card max, **ambient sound only**. We don’t pivot this pre-day-90. The school’s footage becomes ET-recognisable purely through this frame + the look below — without ever overwriting their craft with our branding.

(c) The EducatedTraveler look — a repeatable film identity + two kits

This is what turns a school’s raw phone footage into something that feels like *us* — and it’s deliberately cheap enough for a solo founder at sea to execute, and simple enough to hand to a school. The look serves the Daily Drop’s job: make the most tactile seconds undeniable, and **end on competence**. Distinct from Littoralicious — same warmth, never its CSS or fonts.

The grade, sound, pacing — concrete and repeatable

Element	The ET spec
Capture	Log capture on a recent Pro iPhone (ProRes if storage allows), 4K, 24fps , 180° shutter (1/48). Apps: Blackmagic Camera or the built-in Log capture; lock focus/exposure with AE/AF Lock. (<i>Apple Log availability, app names, and any “free” claim shift release-to-release — confirm the exact device support and app at build time</i>) (Pixflow iPhone guide)
Grade	Log → Rec.709 base LUT, then one fixed ET creative LUT : warm shadows pulled toward ink #0d0b09, highlights held off pure white at the paper tone #f3ede2, a sage-teal lift in mids (#7fa8a5), a terracotta touch in skin/warm objects (#d28a52). Save once, apply to everything for instant consistency (Pixflow mobile 2026)
Sound	Ambient only — the room, the tool, the material. No music bed on the daily clip (also dodges music-licensing risk entirely). One clean diegetic sound carries the whole piece

Element	The ET spec
Pacing	Slow-in, no rapid cuts. 3–6 shots max. Hold the tactile beat 2–3s longer than feels comfortable. Build to one thing
The “end on competence” beat	Last shot is always the finished object or the master’s hands completing the move, held still for a full beat before cut to black. This is the signature emotional close — proof, not promise
The signature open	First 1.5s is the most tactile frame of the whole clip (the hook), <i>then</i> the place-name card fades in over the second shot — never a cold title-first open
Title cards	Display in Fraunces , labels/eyebrows in IBM Plex Mono (the ET fonts — note: distinct from Littoralicious). Warm-dark palette only: paper #f3ede2 text on ink #odob09, sage-teal or terracotta hairline. One card per clip: “{Place} · {Discipline}.” Lower-third credit “Footage: @school” in Plex Mono
Aspect	9:16 vertical, same clip same day to TikTok + Reels + Shorts on @educatedtraveler.app

Edit on a multi-track mobile editor with grading (LumaFusion is a common choice on iOS) so the whole pipeline lives on the phone — pick whichever pro mobile editor you can verify still does the job ([App Whisperer 2026](#)). Build the ET LUT and a Fraunces/Plex title template *once*; after that every clip is a 20-minute job.

Founder-shootable production kit (low-cost, sea-ready)

Prices below are approximate 2026 ballparks — verify current pricing and the exact model before buying; treat the brand names as examples, not endorsements:

- iPhone Pro (already owned) · a compact gimbal (around €100, e.g. a DJI Osmo Mobile) for the slow tactile push-ins · a clip-on ND filter set for the 180° shutter in daylight (around €40) · a small wireless lav for clean ambient/room tone (around €50) · a folding phone tripod (around €20). Roughly €250 all-in, give or take. The iPhone’s lightness is the point — it’s built for observational, over-the-shoulder craft filming ([Pixflow mobile 2026](#)).

Remote / UGC kit for schools to shoot for us

A one-page PDF (Studio-generated, ET-branded) we send at rung 5. Plain instructions, no jargon:

1. **Shoot vertical (hold the phone upright), 4K if you can.**
2. **Three shots:** (a) close on the hands + material, (b) a wider shot of the workspace, (c) the finished piece held still.
3. **No talking, no music** — just let us hear the room.
4. **Daylight or a window; wipe the lens; don’t zoom — step closer.**
5. **Send the raw files (not a screen-recording) via the upload link.**
6. **One line back:** the place, the craft, and “yes you can post and edit these.”

We handle grade, cut, titles, and the ET look. They get the finished film free.

(d) The trust contract — honest, never pay-to-rank, seeds Phase 2 quietly

This is the moat. The instant a school suspects we sell ranking or that a “feature” is a paid placement, the Atlas is worth nothing. So we publish our own rules and live by them.

The four honest commitments (state these openly when asked, and on a public “How the Atlas works” note)

1. **The Atlas ranks by community, never by money.** Position reflects the strength of the scene that gathers there — your best chance of finding your cohort — not reputation, and never who pays. **No one can buy a spot, a rank, or a feature.** (Community-rank, not reputation-rank — Honesty Lock.)
2. **Every card is hand-verified against your official source.** If we can’t verify, we mark it “provisional, verifying” and say so — we never present a guess as fact.
3. **A feature is not an endorsement of a transaction.** Featuring your clip means “we admire the craft and want people to find you.” It is *not* us vouching that any specific course is right for any specific person, and it carries no commercial tie. We say this plainly so a feature is never mistaken for a paid partnership.
4. **We will never surprise you with a sales pitch.** We don’t sell on the site, we don’t take a cut of anything you do, and if we ever (much later) propose hosting a small group together, it’ll be an open invitation, off-site, that you’re free to decline.

How a feature differs from an endorsement — the line we hold

- **Feature** = “look at this beautiful craft, here’s where it lives” → distribution + the Atlas link. Available to any verified school. No money, ever.
- **Endorsement / co-host** = “we’ll personally introduce our community to *you* for a group” → only the Phase-2 shortlist, only after the 8 gates, only off-site. Rare, earned, named for what it is.

Keeping these two visibly separate is what lets us feature 835 schools generously *and* keep the eventual co-host invitation credible.

How this seeds the Phase-2 shortlist — without selling

We don’t recruit partners. We *watch* who climbs the ladder on their own, and log it quietly in the Studio Outreach tab. A school earns a place on the silent shortlist by demonstrating, unprompted, that it’s a real partner:

- replied warmly and gave clean footage rights (rung 3),
- reshared our cut to their own people (rung 4),
- showed up well in a co-created shoot (rung 5),
- and — the strongest signal — **our clip of them genuinely travelled** (a Phase-2 gate in itself), or their craft drew unprompted “can you connect me?” inbound to the Circle.

None of this is ever communicated as “you’re being evaluated as a partner.” It’s just relationship-building. The shortlist is a private byproduct. When the 8 Phase-2 gates finally turn green together, we already have a ranked, trust-tested list of one or two schools who’ve shown they’re real — and we open rung 6 to them first.

The discipline that protects all of it: every rung 1–5 move must be able to stand alone as a gift. The day a school can tell we were only being generous to set up a sale, the moat is gone. Generosity now is what makes the

Phase-2 ask land later — and what keeps the Circle filling with people who open, click, and refer between now and then.

Sources (third-party 2026 analyses and guidance pages — directional, not authority; verify before betting on them): Showcase — UGC usage rights 2026 · Later — repost rules & credit · Digital Applied — UGC licensing framework 2026 · The Social Media Law Firm — UGC legal risks 2026 · Automateed — reposting rules · Pixflow — cinematic iPhone 2026 · Pixflow — mobile filmmaking 2026 · App Whisperer — best iPhone filmmaking apps 2026

Part 4 — The Community: a social home for skills

Here is the one thing to hold onto before any feature: **we are not building a feed. We are building a room where people who chase the same fire can show each other their hands at work.** The feed is just the wall they pin things to. Everything below bends toward the one metric — engaged Circle subscribers, people who open, click, reply, and refer — and away from the vanity the rest of social chases. A thousand people who truly care beats a million passing through, and the research on niche communities in 2026 says the same thing plainly: small and active beats big and empty ([mean.ceo](#)).

So the test for every post, prompt, and pixel in here is the master test from the Focus Filter: **does it fight the screen or feed it?** A skills-only social home that turns into another dopamine treadmill has betrayed its own reason to exist.

Honesty caveat for this whole section: the external numbers and platform/policy claims below come from third-party 2026 analyses and help-center pages. They are directional, not gospel — verify them in the Studio Metrics tab before you bet a format on them. Our moat is honesty; that discipline starts inside our own guideline.

(a) Vision + the strictly phased MVP

The vision, in one line: *The Circle grows from a list you read into a room you post in — a place, a person, your people, made visible.* Not “social media for everyone about everything,” but the opposite: social media for one person at a time chasing one real craft, anchored to one real place on the Atlas.

We already own most of the foundation. The job is to **deepen what’s live, in order**, never to leap to the full feed. Here is the build, phased so each step proves engagement before the next is allowed.

Phase	What ships	Built on (live today)	The bar to advance
0 — Now (no new feed)	Turn cohort-chat + crew cards into a <i>reason to post</i> . Add a weekly founder-led ritual thread (“This week’s hands”). Crew cards get a single “what I’m working on right now” line.	cohort-chat (realtime group chat), crew cards, community sidebar, Circle list	People reply to the ritual thread unprompted, 3 weeks running
1 — The Pinboard (MVP feed)	A single, slow, per-cohort wall where members pin one	cohort-chat tables + Supabase profiles	A cohort wall stays alive without the founder posting

Phase	What ships	Built on (live today)	The bar to advance
	skill post. No global feed yet. Chronological inside a cohort, no algorithm.		daily
2 — Per-craft + per-core rooms	The wall splits by discipline and then rolls up into the four cores. Atlas anchors every post. Surfacing by community-rank.	Atlas (99 disciplines / 4 cores), discovery lenses	Cross-cohort replies happen; a clip “travels” between rooms
3 — Bridges out	Feed → IRL (who’s-going / local circles) and feed → Skill Passport (progress posts feed prior-experiences + instructor-verified flag).	prior-experiences table, instructor-VERIFIED flag, IRL mechanism (build target)	Phase-2 GATE territory — only once the 8 gates approach true

What we explicitly do NOT build now (Stop-Doing): no global cross-everything feed, no follower counts, no like-leaderboard, no infinite scroll, no recommendation algorithm, no re-surfacing of selling. The Atlas is done at 99/424 — distribution is the work, not more pages.

One-idea-per-phase rule: each phase exists to move engaged subscribers, not to add surface. If a phase doesn’t raise opens/replies/referrals, we don’t advance — we fix it.

(b) The feed model — what a “skill post” is

A skill post is **not a status update**. It is the smallest unit of “I did a real thing with my hands today.” Four post types, and only four, so the room never drifts into chatter:

1. **Progress** — a short clip or before→competence image of you actually doing the craft. (“First clean roll-tack. Took me forty tries.”)
2. **A hard-won fact** — one concrete thing you learned at the source, the kind no video gives you. (“The miso changes flavour by the hour in summer — the master taps the lid, doesn’t open it.”)
3. **A beautiful image** — the material, the tool, the place, the finished object. Tactile, not staged.
4. **A real question** — pointed, answerable by a peer or a master. (“Why does my knife drag on the pull and not the push?”)

Every post must anchor to the Atlas. This is the load-bearing rule and our entire moat against the skillcation-listicle swamp. A post cannot be published without selecting: - a **discipline** (one of 99), which sets the **core** colour (wellness #94ad86 / adventure #6fa3a0 / creative #cf8f6e / culinary #c9a24a), and - a **place** (an Atlas destination — the source, a school, or “where I am working it”).

That single requirement does three things at once: it ties the social home to the real map instead of the void, it routes the post into the right per-craft room, and it makes every post a quiet ambassador for an Atlas page (which serves the one metric — engaged subs find real masters).

The structure — per-craft, rolling up to per-core:

CORE: Culinary (#c9a24a)

└ Discipline room: Sushi – Tokyo / Kanazawa

└ skill posts, chronological, anchored to Atlas place

└ Discipline room: Bread – San Sebastián

CORE: Adventure (#6fa3a0)

└ Discipline room: Big-boat sailing – Mediterranean

A member lives mostly in their **discipline room** (their people). The **core view** is the wider neighbourhood they can wander into. There is no “everything” view — that absence is a feature, not a gap.

Surfacing — by community-rank, never popularity (Honesty Lock). The Pinboard orders posts by: - **Recency within your discipline room** first (you see your people, now). - **Community-rank of the place** second — the same five-dot “best chance to find your cohort” signal the Atlas already uses, NOT a like-count or reputation score. - **“A peer replied” boost** third — a post that drew a real answer rises, because a conversation is the thing we want more of.

No engagement score. No “trending.” No view counts shown. The question we ask of the ordering is the same as everywhere: does this pull someone toward *doing*, or toward scrolling?

(c) Mechanics that keep it authentic, not generic social

The 2026 evidence is blunt: infinite scroll and algorithmic feeds work like slot machines — intermittent, uncertain reward firing the dopamine loop — and the EU has signalled it intends to act on dark patterns and addictive design (a “Digital Fairness Act” has been floated alongside a fitness-check on existing rules; treat the exact name, status and effect as provisional and verify before quoting it) ([Jumpstart](#), [The Week](#)). We build to that standard now, by choice, not because a regulator made us. Calm-technology principle: the room should sit at the periphery and reward leaving it, not trap attention at the centre ([Jumpstart](#)).

Concrete mechanics:

- **Finite, not infinite.** The Pinboard loads one cohort’s posts, then *ends*. A clear “you’re caught up — go make something” card at the bottom. No auto-load, no “load more” pulling you down a hole.
- **Quality gate over reach.** Posts that aren’t doing-the-craft (selfies, links, hot takes, “good morning”) simply have no post type to live in — the four-type structure refuses them by design.
- **No follower counts, no like tallies, ever.** A post can be marked **“this helped me”** by a peer — but the count is private to the author, never a public number. Recognition, not a scoreboard.
- **Real-name, real-progress culture.** Profiles carry real names and a real craft. The crew card’s “what I’m working on right now” line is the social object — progress, not persona.
- **Prompts that pull toward DOING.** The weekly ritual is a *task*, not a topic: “Post your worst attempt this week.” “One fact your master gave you that a video never would.” “Show the tool, not the trophy.” This is straight from what works in small communities — rituals that lower the effort to participate and ask for *specific* contributions, not generic intros ([innoloft](#), [mean.ceo](#)).
- **No notifications-as-bait.** One digest, not a stream of pings. A weekly “in your room this week” email (rides the existing Circle/Resend sender) — which doubles as direct fuel for the one metric: opens and clicks.
- **Slowness is the brand.** Daily Drop stays the wordless tactile clip; the feed inherits its DNA — hands, material, the moment of competence — never talking-heads or text walls.

*The line to hold: **quality over reach, the source over the screen, your people over an audience.** Everything cold, glassy, gamified, or Netflix-shaped is retired forever.*

(d) Moderation, safety, trust at solo-founder scale

Arnaud is solo, part-time, often at sea. Moderation must be lightweight, honest, and defensible — and the literature is clear that trust infrastructure has to exist *before* the first incident, not be bolted on after ([mean.ceo](#)). The good news: a small, anchored, real-name room is most of the safety by construction.

Structural safety (does the work without the founder): - **Members only.** Posting requires a Supabase account in a cohort — no anonymous drive-by, which kills most abuse before it starts. - **Real name + real craft + Atlas anchor.** Three friction points that filter out spam and bad actors without a moderation team. - **Small rooms.** Per-discipline scale means peers self-correct; the founder is a presence, not a police force.

Lightweight human layer: - **One page of house rules, in the founder voice.** Short: this is a room for people doing real things; be kind, be specific, anchor to a real place; no selling, no spam, no cruelty. Visible at the door (set the tone early — clear rules, visible moderation, respectful-communication guidance ([bettermode](#))). - **Report → quiet queue.** A single “flag” on any post drops it into a Studio queue (build on the existing [studio.html](#) cockpit). Flagged posts auto-soften (hidden from the room, visible to author) pending a founder glance — fail-safe, not fail-open. - **Rate limits, not algorithms.** A few posts per member per day. Stops flooding without judging content. - **Honest absence.** If the founder is at sea for a week, the room is told: “Arnaud’s offshore — slower replies, same room.” Honesty is the moat; we never fake presence or counts.

Trust mechanics tied to the Honesty Lock: - Self-declared skills/credentials stay clearly *self-declared* (current free-text reality). - Only the **instructor-VERIFIED flag** on prior-experiences carries weight, and it’s marked as such — never inflated into a fake badge. - No fabricated member counts, no “10k+ makers” puffery. We show what’s real.

(e) Seeding, cold-start, and the bridges out

Cold-start is where most communities die quietly — and the fix is well-documented: seed before you launch, invite a narrow first cohort, ask for specific contributions, and expect 6–12 months to real life ([innoloft](#), [mean.ceo](#)). Apply it literally:

Seeding (now-move): 1. **Founder goes first, daily, for weeks.** Arnaud posts real progress from his own crafts (cooking, sailing) anchored to real Atlas places. The room must be *full of doing* before anyone is invited in — empty rooms broadcast death. 2. **Narrow first cohort.** Pull the most engaged existing Circle subs (the openers/clickers — they’re already identified in Resend/analytics). Invite by hand, by name, into one discipline room. Not a broadcast. 3. **Specific asks, not “introduce yourself.”** First prompt: “Post one thing you made this month and the one place you’d go to get better at it.” This produces Atlas-anchored posts on day one. 4. **Instructor cross-post (now-move, no money).** Use the existing invite-instructor edge function. A verified school/master shares a single learning clip or fact into the relevant discipline room. This is the spine made real — *a person* shows up where *your people* are — and it warms the room with authority, not ads.

Bridge: feed → IRL (phase-gated later-move). A “who’s-going / local circles” mechanism on the discipline room: when a place lights up with posts, members can quietly raise a hand for an in-person gathering at the

source. This is the IRL connection mechanism that’s a build target — and it only matters once rooms are alive (Phase 3 / approaching the 8-gate).

Bridge: feed → Skill Passport. A **progress** post can be promoted by its author into their **prior-experiences** record. If a master/instructor saw it, the instructor-VERIFIED flag can attach. The feed becomes the *evidence stream* that feeds the credential — the CREDENTIAL + CURATION + COHORT fusion no incumbent stacks. Honest, earned, not bought.

Both bridges feed the one metric in opposite directions: IRL gives the “can you connect me?” inbound the Phase-2 gate watches for; the Passport gives subscribers a reason to stay engaged for months, not days.

(f) Honest anti-patterns we refuse

Tied to the Stop-Doing list and the Honesty Lock — written down so they’re never quietly re-introduced when growth feels slow:

We refuse	Why	What we do instead
Infinite scroll / autoplay	Slot-machine dopamine loop (Jumpstart)	Finite cohort wall, “go make something” end card
Follower / like counts as public scores	Vanity; not the one metric	Private “this helped me,” real-name presence
Algorithmic ranking / “trending”	Manufactures outrage + addiction	Recency + community-rank, never popularity
Engagement-bait prompts (“tag 3 friends”)	Reach over substance	Doing-prompts (“post your worst attempt”)
Fake member counts / inflated badges	Breaks the Honesty Lock (the moat)	Show only what’s real; mark self-declared vs verified
Notification streams	Attention farming	One weekly digest (also serves opens/clicks)
Re-surfacing selling / prices / checkout	Strategy Lock — community first	Money is Phase-2, off-site, after 8 gates
A global “everything” feed	Generic social swamp	Per-craft rooms, Atlas-anchored

The discipline here is the product. Anyone can build a feed; the rare thing is a social home that refuses the cheap tricks and still feels alive. That refusal — angle, voice, focus, media muscle — is the moat the funded incumbents can’t copy, because they’re built on the metrics we’re walking away from.

The close, one thing: build the room slowly, fill it with real hands first, anchor every post to a real place, and judge it only by whether the Circle opens, replies, and refers. Skills last, tans fade — and a social home for skills only earns the name if every screen in it points back to the source.

Sources (third-party 2026 analyses — directional, verify in Studio Metrics): [mean.ceo — niche community o to 1,000](#) · [mean.ceo — customer community guide](#) · [Jumpstart — dopamine loop](#) · [The Week — end of infinite scroll](#) · [innoloft — how to build community 2026](#) · [bettermode — grow a community 2026](#)

Part 5 — The Skill Passport: profile + credentials

A crew card is a snapshot. A passport is a record. The move here is to take the read-only crew card you already ship at `/crew` and turn it into one public page a person is proud to put in their bio, in their email signature, on their CV — a page that says *here is what I can actually do, here is who taught me, and here is the proof*. Skills last, tans fade. The passport is where that line stops being a slogan and becomes a document.

One discipline before anything else: **the passport’s only job is to make the honesty visible**. Every other publication on earth lets you type “expert” into a box. We don’t. The instant a stranger can’t trust a line on the page, the page is worthless — and so is the moat. So we build the profile first, and we build the credential model second and slower, and we refuse to issue a single thing we can’t stand behind. That refusal is the product.

Caveat on the figures below. *The credentialing-standard dates, the LinkedIn integration behaviour, and the “what employers value” claims in this section are drawn from third-party 2026 analyses and provider docs — directional, not gospel. Verify each one in the Studio Metrics tab (or against the primary source) before betting the build on it. For a brand whose moat is honesty, we never restate someone else’s blog stat as settled fact on a public surface.*

(a) The public profile — link-in-bio that earns the next click

Today `/crew` is a read-only card the visitor lands on cold. The passport is the same data, reframed as a **landing page a holder actively shares** — a Linktree that happens to be a verifiable record.

Share URL. Keep it human and ownable. The holder picks a handle at the profile; the page lives at:

educatedtraveler.app/p/yann-le-bihan

A vanity slug (not the Supabase UUID) is the single thing that makes a page *shareable* — nobody puts `?id=8f3a-...` in their bio. Add a slug column to `profiles`, unique, claimed once. The QR posters you already cut (`marketing/qr/qr-crew-*`) point here. The existing `/crew?id=` route 301s to `/p/<slug>` so nothing breaks.

What the page shows, top to bottom (one idea per block — never lose focus):

Block	Content	Why it’s there
Header	Name, one-line “what I’m chasing” (e.g. “ <i>learning to roll a knife in Tokyo</i> ”), home city, avatar, the four-CORE accent of their main craft	The hook — a person, not a résumé
The Skills	<code>skills[]</code> rendered as chips, each carrying its tier glyph (see below)	The payoff, above the fold
The Record	<code>prior_experiences</code> rows — <i>what, with whom, where, when</i> + the Verified-by-instructor mark where the flag is true	This is the part no competitor has
The Links	The link-in-bio strip: their Daily-Drop clips, their public posts in the skills feed,	The funnel — every link goes <i>inward</i>

Block	Content	Why it's there
	their Atlas pages for the crafts they chase, “message me on the Circle”	
The footer	“On the Atlas at [discipline] · [place]” + one quiet line: “A place, a person, your people.” + the join-the-Circle field	The re-hook to THE ONE METRIC

The funnel — every link points home. A Linktree usually scatters people *outward* to a dozen platforms. Ours does the opposite: a passport’s links are the holder’s own Daily-Drop clips (which sit on our feed), the Atlas pages for their craft (our owned SEO), and a **“connect with me” button that drops the viewer into the Circle join flow** (`#circle` → `launch_waitlist`, the front door you already deployed). A visitor who lands on a friend’s passport, sees real proof, and clicks “I want this too” becomes an engaged Circle subscriber. That is the whole game: **the passport turns one member into a recruiting surface for the next.** Word-of-mouth is the cheapest channel a solo founder has, and the passport is the artifact that carries it.

How it ties to the three existing surfaces: - **Feed** → the passport’s “Links” pulls the holder’s own posts/clips from the skills-only feed; a skill chip is clickable into the feed filtered to that craft. - **Atlas** → each skill links to its Atlas discipline page; each `prior_experiences` location links to its Atlas destination. The passport *feeds the Atlas its inbound links* and the Atlas gives the passport its credibility (“this place is real, here it is on the map”). - **Circle** → the connect button and footer field route to the Circle. Visibility stays governed by the column you already have: `private` / `cohort` / `public` — the passport is simply the rendered `public` view, opt-in, never default-on.

(b) The credential model — four tiers, brutally honest

The current state is the right foundation: `profiles.credentials` is free text, `profiles.existing_certs` is free text, and only `prior_experiences.verified` carries real weight (an instructor flips it on a row they can vouch for). We keep that exact mechanic and wrap a **tier ladder** around it. The tier is not decoration — it changes how the claim is *displayed and described* to a stranger.

Tier	Name	Source of truth	How it shows	What it honestly means
0	Self-declared	The holder typed it (<code>skills[]</code> , <code>credentials</code> text)	Plain chip, label “ <i>Self-declared</i> ”	“This person says they can do this.” Nothing more.
1	Community-vouched	Peers in the holder’s cohort confirm it	Chip + small count, “ <i>Vouched by 3 in the circle</i> ”	Social signal, not proof. Honest <i>because</i> we label it as soft.
2	Instructor-verified	The teaching master / school flips the existing <code>verified</code> flag	Chip + ✓ Verified by [Instructor], [School] + the Atlas link	“The person who taught them confirms they were there and did the work.” This is the strong, real tier.

Tier	Name	Source of truth	How it shows	What it honestly means
3	ET-issued	EducatedTraveler signs a credential	✓✓ EducatedTraveler • Completion of [curated immersion]	The <i>only</i> thing ET issues — and only for what ET can legitimately attest. Phase-2-gated.

The honesty rules — read these twice, because breaking one breaks the moat:

1. **ET can certify ATTENDANCE and PARTICIPATION at things ET ran. ET cannot certify MASTERY.** A Tier 3 record would read, hypothetically, “*completed the ET [craft] immersion in [place], [dates], with [master]*” — a record of presence and engagement at a curated ET immersion. (No such immersion exists or is scheduled today; this is an illustrative shape only — Tier 3 is hard-gated, see the phasing below.) It does **not** say “*certified sushi chef.*” We did not assess a skill against a standard; we hosted a week. Say exactly that and no more.
2. **The teacher owns mastery, not us.** If anyone gets to attest “this person can actually do the thing,” it is the master who watched them do it — Tier 2, the instructor flag. ET’s job is curation and connection, not examination. This maps to the founder’s real moat (media/curation) and stays inside the strategy LOCK: *introduce / connect / bridge*, never assess-and-certify like a credentialing body we are not.
3. **Never let a lower tier wear a higher tier’s clothes.** A self-declared “expert woodworker” must look visibly different from an instructor-verified one. The tier glyph and label are mandatory and non-removable. The day a Tier 0 chip is indistinguishable from a Tier 2 chip, we are a résumé-inflation site, and we have become exactly the “*certificate from an unknown online platform with no industry recognition*” that — per 2026 third-party analyses (directional, verify in Studio Metrics) — employers report discounting (Sertifier, 2026).
4. **No issuance for watching.** One 2026 hiring write-up argues employers do not value “*credentials that require nothing more than watching videos and paying a fee*” (The Interview Guys, 2026) — treat that as a directional claim to validate, not a settled law, but it points the right way. Tier 3 is gated behind a real, embodied, week+, at-the-source immersion that passes the FOCUS FILTER. If it doesn’t return a real skill — embodied, immersive, community-based — it cannot earn an ET credential. The Focus Filter is the credential’s admissions test.
5. **Revocation is a feature, not an embarrassment.** A teacher can un-verify. ET can revoke a Tier 3. A verifiable credential the holder can’t lose is just marketing; one that can be revoked is trustworthy. We build revocation in from day one (see standards below).

The credential-inflation trap, named so we never walk into it: the temptation, the moment a few hundred people hold passports, will be to mint “EducatedTraveler Certified ____” because it looks impressive and people will ask for it. Resist completely. The instant ET attests to a skill it never assessed, the verified mark means nothing, and the one thing we have that VAWAA, MasterClass and the skillcation-listicle swamp don’t — *trustable proof* — evaporates. **We would rather issue ten credentials a year that are bulletproof than a thousand that a hiring manager rolls their eyes at.**

(c) The standards — portable, provable, low-ops

Three things appear true in 2026 and worth building on (each is provider/standards-body documentation — confirm the exact status and dates at build time, since these move):

- **Open Badges 3.0** was reported voted Final by 1EdTech in mid-2024 and is now widely treated as the dominant skill-credential standard, with cryptographic signing built in ([1EdTech impl guide](#); [virtualbadge.io, 2026](#)). Verify the “Final” status with 1EdTech before relying on it.
- **Open Badges 3.0 is itself a W3C Verifiable Credential.** It is designed to align with the W3C Verifiable Credentials Data Model v2.0, which is reported to have reached W3C Recommendation in mid-May 2025 — same `issuer` / `credentialSubject` / `achievement` / `proof` / `credentialStatus` shape ([Sertifier](#); [VC Data Model status](#)). Confirm the exact Recommendation date against W3C at build time. If it holds, adopting OB 3.0 gets us VC compliance for free — one standard, two boxes ticked.
- **Trust comes from issuer identity, evidence, and stable identifiers — not from a badge image.** The 2026 field guidance is explicit that a PDF or a pretty PNG is worthless; verifiers check *who signed it, what the evidence was, and whether the identifier resolves* ([Sertifier credentials guide, 2026](#)).

What that means for a Tier-3 ET credential (Phase 2): issue it as an **Open Badges 3.0 credential**, which gives us, in one JSON object:

- `issuer` → ET’s identity, anchored with `did:web` keyed to `educatedtraveler.app`. This is the low-ops choice: a `did:web` is just a signed JSON file served from your own domain — no blockchain, no fees, no third party. The conformance guidance even recommends `did:web` on the issuer’s own web domain precisely *to avoid portability challenges* ([UN Transparency Protocol VC spec](#)).
- `credentialSubject.achievement` → the honest claim (“Completed the EducatedTraveler [craft] immersion, [place], [dates], with [master/school]”).
- **evidence** → an `evidence` URL pointing back to the holder’s passport / the Atlas page / the Daily-Drop clip of their work. Evidence is what separates a real credential from a participation sticker.
- **proof** → ET’s cryptographic signature. Sign once at issuance with the `did:web` key.
- **credentialStatus** → a **Bitstring Status List** (the W3C revocation mechanism) — a single small file at `educatedtraveler.app/credentials/status/1` that lists which credentials are still valid. To revoke one, flip one bit. It’s privacy-preserving and storage-agnostic; the issuer just serves it from a URL ([W3C Bitstring Status List](#); [walt.id status docs](#)).

The concrete low-ops path for a solo founder (do NOT build a credentialing platform):

1. **Don’t self-build crypto.** For the handful of Tier-3 credentials Phase 2 will produce, **use a managed OB 3.0 issuer** (Accredible, Canvas Credentials/Badgr, Credly, POK — reported as certified 3.0 issuers in 2026; confirm certification before choosing) OR a small open-source signer (`walt.id`, `transmute status-list.vc`). A solo founder issuing dozens, not millions, of credentials should not run key infrastructure by hand.
2. **Anchor issuer identity as `did:web`** under your own domain — one static signed file, version-controlled in the repo, deployed via the existing `gh-clone-to-/private/tmp` recipe. No new infra.
3. **Serve the status list as a static file** on the same GitHub Pages site. Revocation = edit one file, redeploy. This fits your stack exactly — it’s the same deploy muscle as the Atlas.
4. **LinkedIn “Add to profile.”** This is the high-leverage integration and it costs nothing: LinkedIn’s *Add to Profile* lets an issuer generate a pre-filled link (name, org, issue date, credential ID, credential URL) that drops the credential straight into a member’s *Licenses & Certifications* section

(addtoprofile.linkedin.com; [LinkedIn Help](#)). Generate one button per Tier-3 credential. As of the latest provider guidance, the public Add-to-Profile link reportedly needs no API approval (the full Profile Edit API is gated to approved partners, which we don't need) — **verify this against LinkedIn's current terms at build time, since their policy changes**. The credential URL points back to the holder's ET passport, so LinkedIn becomes a *funnel back to us*.

Tier 0–2 need none of this. They are not signed credentials — they are honest, labelled claims living in your Supabase rows, rendered on the passport. Don't over-engineer them. Open Badges machinery is reserved for Tier 3, where a stranger needs cryptographic proof an institution stands behind it.

(d) Usefulness — for the holder, for the verifier

For the holder (work / CV): - The passport URL goes in an email signature, a CV, a portfolio, an Instagram bio. It says *"I spent a real week learning a real thing from a real master"* — which, per 2026 third-party trend analyses (directional, verify in Studio Metrics), is the kind of non-degree signal a skills-based hiring market is said to weight more heavily ([virtualbadge.io trends, 2026](#)). - A Tier-2 instructor-verified record is genuinely useful even *without* a Tier-3 cert: "Verified by [named master] at [named school]" is a checkable reference, which is more than most online certificates offer. - A Tier-3 OB 3.0 credential pushes to LinkedIn in two clicks and travels anywhere — it's a portable W3C VC, not a walled-garden badge.

For the verifier (a hiring manager, a school, a charter captain checking a chef): - Tier 0–2: they open the passport, see the tier labels, and click through. An instructor-verified row links to the named instructor and the Atlas school card — they can contact the source. **Transparency is the verification.** - Tier 3: they paste the credential into any W3C VC verifier (or click "verify" on the passport), which resolves ET's `did:web`, checks the signature, and checks the Bitstring Status List for revocation. Green = ET signed this, it hasn't been pulled, here's the evidence URL. This is the *"who signed it, what's the evidence, does the identifier resolve"* test the 2026 guidance demands — and we pass it by construction.

(e) Phasing — ships now vs. Phase-2-gated

Ships NOW (builds only on what's already deployed; advances THE ONE METRIC):

- ☐ **The better profile** — vanity `slug` column on profiles; `/p/<slug>` route; 301 from `/crew?id=`. Warm-Dark Editorial, same tokens as the Atlas.
- ☐ **Tier labelling on existing data** — render `skills[]` as Tier 0 (Self-declared) chips; render `prior_experiences` rows with the Tier-2 ✓ mark wherever `verified = TRUE`. This is *display logic over the live schema* — no new tables.
- ☐ **Tier 1 (community-vouched)** — a lightweight `skill_vouches` table (one peer, one skill, one row, RLS owner-can't-vouch-self). Optional for v1; it earns its place only once cohorts are active.
- ☐ **The link-in-bio strip + connect-to-Circle button** — the funnel. Every link inward; the connect button routes to the Circle join flow.
- ☐ **Instructor-verify UX** — make flipping the `verified` flag a one-tap action in the instructor surface (the policy already exists; just surface it). More verified rows = more trustable passports = more shares = more Circle sign-ups.

Phase-2-GATED (do NOT build until all 8 gates are true — especially a real curated ET immersion has actually run, gate 7):

- ☐ **Tier 3 ET-issued credentials** — `did:web` issuer identity, OB 3.0 issuance via a managed issuer, Bitstring Status List for revocation, LinkedIn Add-to-Profile button. There is *nothing to issue* until an ET immersion has happened, so this is structurally Phase-2 by definition. Building the signing pipeline before there's a single legitimate thing to sign would be exactly the credential-inflation we swore off.

Why this phasing protects the metric and the moat: the now-moves make every existing member's page a recruiting surface that drives engaged Circle sign-ups — the only number that matters in the next 90 days — using data and schema you already have. The Phase-2 move waits, correctly, because a credential issued before there's something real to certify is the fastest way to turn the honesty moat into a liability. Earned, not bought — including the right to issue the credential itself.

Sources (all third-party 2026 analyses or standards-body / provider docs — directional, verify before relying on any specific figure, date, or policy claim): - Open Badges 3.0: What Is the Status in 2026? — virtualbadge.io - Open Badges 3.0 explained — Sertifier - Open Badges 3.0 Implementation Guide — 1EdTech - Credentials in 2026: trust & verification — Sertifier - Micro-credentials that actually get you hired — The Interview Guys - Top 5 Digital Credentialing Trends 2026 — virtualbadge.io - W3C VC Bitstring Status List (revocation) - Credential Status & Revocation — walt.id - Verifiable Credentials / `did:web` — UN Transparency Protocol - LinkedIn Add to Profile · Manage Licenses & Certifications — LinkedIn Help

Part 6 — IRL: connecting like-minded people

The whole point was never the screen

Everything we build online — the Atlas, the Daily Drop, the Circle — is a means, not the end. The end is two people who were chasing the same fire finding each other, then standing in the same room with their hands in the same clay. “A place, a person, your people.” Parts 1 through 5 earned the attention. This Part spends it: it turns a click into a coffee, a coffee into a practice circle, and — once the 8-gate opens — a practice circle into a week at the source.

One rule governs every rung below: **IRL is not a new product. It is the deepest expression of the one metric.** A member who shows up in person, meets a peer, and comes home glad will open the next email, click the next Atlas page, and tell one friend. So every rung here is judged the same way as a post or an article — does it produce **engaged Circle subscribers**, and does it pass the **focus filter** (embodied, immersive, community-based, returns a real skill)? If a rung feeds the screen instead of fighting it, it's out.

And the honest constraint that shapes all of it: Arnaud is solo, part-time, often at sea. **The founder cannot be in the room.** So the entire IRL design is built to run *without* him — light enough that a volunteer host carries it, safe enough that strangers can meet without ET becoming an operator, and instrumented so the proof flows back into the feed on its own.

***On the comparison figures below:** every external number in this Part (MEA, AI Tinkerers, meetup programs) comes from third-party 2026 pages and is directional, not gospel — a pattern to learn from, not a metric to bet the format on. Verify any figure in the Studio Metrics tab before you lean on it. Our own numbers are the only ones we treat as fact, and our honesty lock applies here as much as it does to Atlas facts.*

(a) The offline ladder — lightest to deepest

Five rungs. Each is a smaller ask than the next, and each one a member climbs is a stronger Circle subscriber. The first four are **now-moves** (they need no money, no booking, no operations). The fifth is **Phase-2-gated** and stays dark on the public site until all 8 gates are green.

Rung	What it is	The ask	Founder ops	Gate
1. Near me	“Who else near me is into X” — a soft signal, never a directory	Tick a box	Zero (passive)	NOW
2. Who’s going	A “who’s going / who’s been” marker on Atlas places	Tap a heart	Zero (passive)	NOW
3. Practice circle	Informal local skill meetup — members run it themselves	Show up	Near-zero (template only)	NOW (soft, opt-in seed)
4. ET gathering	A small ET-hosted in-person evening	Show up	Light (1 event, founder-presence optional)	NOW (occasional, when capacity allows)
5. The immersion	The cohort week at the source, co-hosted with a vetted partner	Travel together	Heavy — Phase-2 only	GATED (8-gate)

Rung 1 — “Who else near me is into X.” Inside `community.html`, a member can opt in to a single, coarse signal: *city or metro region* (never precise location) + their top one or two crafts. The product never shows a list of named strangers. It shows a **count and a nudge**: “3 others near Lisbon are chasing wood-fired ceramics. Want us to tell you when there are enough to gather?” That’s it. The count is the hook; the connection is consent-gated (see (c)). This rung costs nothing to run and converts passive readers into members who’ve raised their hand — the exact behaviour the one metric measures. It builds on what’s already live: the profile already stores `interests` and `skills[]`, so this is a query and a toggle, not a new system.

Rung 2 — “Who’s going / who’s been” on Atlas places. The Atlas is 324 destinations of cold facts today. We warm one corner of it with a single peer signal. On a destination or school card, a member can mark — privately, opt-in — “*I’ve been here*” or “*I want to go*.” The card then shows an aggregate, never names: “4 from the Circle have trained here · 11 want to go.” This is the lightest possible expression of “your people,” and it’s pure focus-filter fuel — it tells a reader that a real human peer walked through that door. It also feeds the community-RANK honesty (best chance to find your cohort), not a reputation score. Crucially, it does **not** become a booking widget or re-surface selling — it’s a heartbeat, not a checkout. (This is the *only* new thing we add to the Atlas; the stop-doing list forbids building more Atlas pages, and we honour it.)

Rung 3 — Practice circles (local skill meetups). When a city crosses a small threshold of members chasing the same craft (start the threshold honestly low — 3 to 5), we don’t run an event. We *invite a host*. A practice circle is a recurring, member-run gathering: a few people who knead, throw, sail, sketch, or cook *together*, monthly, in a real place a member already has (a studio, a kitchen, a beach, a park). ET supplies the spark, the template, and the introduction — nothing more. This is the load-bearing rung: it’s where “community-based” stops being a slogan. The volunteer playbook for it is in (b).

Rung 4 — ET-hosted small gatherings. Occasionally — when Arnaud is ashore in a city with a real cluster, or a trusted host steps up — ET puts its name on one evening: a screening of Daily Drops on a wall, a hands-on demo from a local master we’ve Atlas-verified, a long table. Small (8–20), warm, no selling, no pitch. The point is to let members feel the brand as a *room*, not a feed, and to recruit practice-circle hosts from the people who light up that night. Founder presence is **optional** — a host can run it from the same template. We keep these rare on purpose; rarity protects founder capacity and makes each one matter.

Rung 5 — The cohort immersion (GATED). This is the destination of the whole ladder: a week+ at the source, with a vetted partner master, alongside the peers a member already met on rungs 1–4. It is **off-site, referral/co-host fee only, never an on-site transaction**, and it stays invisible on the public site until every one of the 8 gates is true — including gate (4), *regular unprompted “can you connect me?” inbound*, which the lower rungs are designed to manufacture. We do not soften this. Until the gate opens, the answer to “how do I do the immersion?” is “join a practice circle and come to a gathering — that’s where it starts.” Honest, and it drives the one metric instead of revenue we can’t yet serve.

The ladder’s logic: each rung is a smaller, warmer ask that produces a stronger signal of intent. By the time someone has been to three practice circles, the immersion isn’t a sale — it’s the obvious next step among friends. That’s “earned, not bought,” made literal.

(b) Local Circles / chapters — self-organising with minimal founder ops

The model to adapt — honestly — is **MEA’s**: per MEA’s own pages, dozens of volunteer-run regional chapters (figures vary by source and are approximate) across a reported ~60 countries serving a reported 8,000+ alumni, with chapter events posted on a shared portal calendar and members meeting for dinners, hikes, talks, and structured conversations entirely on their own steam ([MEA Wisdom Community](#), [MEA alumni portal](#)). The same volunteer-host pattern reportedly runs AI Tinkerers’ large multi-city network and the WordPress and FSF meetup programs — events “organized by any reliable, trusted member of the community” off a provided guide and a safe-space policy template ([WordPress Meetup Program Basics](#), [FSF LibreLocal](#)). The lesson across all of them — and the part we actually borrow — is structural, not the headcount: **the founder supplies a kit and an introduction; volunteers supply the room**. That is the only way a solo, at-sea founder runs a global IRL community.

But we adapt it honestly. MEA chapters are bound by a shared *workshop experience* members already paid for. We have no shared experience yet — so a Circle chapter is bound by a shared **craft**, not a shared alumni status. A chapter is “Lisbon · Ceramics,” not “Lisbon.” That keeps the focus filter intact: every chapter exists to practise one real thing together.

How a chapter is born (seeded by online activity): 1. Rung-1 signals cross threshold in a city for one craft (e.g. 5 members near Lisbon ticked ceramics). 2. The Studio cockpit surfaces it on a new “**Sparks**” view (Metrics-tab sibling) — a ranked list of city × craft clusters ripe to seed. 3. ET emails the cluster one message: “*There are now enough of you near Lisbon chasing wood-fired ceramics to gather. Would one of you host the first practice circle? Here’s everything you need.*” 4. First volunteer to say yes becomes the **host**. ET introduces them to the others (consent-gated, see (c)) and hands over the playbook. Done. No founder presence required.

The volunteer host playbook (one page, lives at [community.html](#) → “Host a circle”): - **What a circle is:** 3–8 people, one craft, one recurring meet (aim monthly), in a real place *you already have access to* — your studio, a kitchen, a beach, a park, a community room. ET does not rent space or pay for anything. - **Your job**

(small): pick a date, pick the place, tell people what to bring, show up, make it warm. That's all. You are a host, not a teacher and not a guide. - **Your first message template** (ET supplies it pre-written): time, place, what to bring, what we'll make/practise, "all levels welcome, we're here to do the thing together." - **The four house rules** (printed on the kit, non-negotiable): (1) we make/practise something real every time — no talking-shop-only meets; (2) all levels welcome; (3) the safe-space policy applies (see (e)); (4) no selling, no recruiting, no MLM — this is a craft circle, not a pitch. - **After each meet:** post one photo or 30-second clip to the chapter thread and tag it #FromTheCircle. (This is the loop-back engine — see (f). It's the *only* thing ET asks back.) - **What ET gives you:** the intro to your members, the message templates, the safe-space policy, a printable one-pager, a spot on the Circle calendar, and a feature in the newsletter when your circle does something good. **What ET does not give you:** money, insurance, space, or instruction. You're the host; ET is the matchmaker.

Minimal-ops guarantees. A chapter must survive Arnaud being at sea for a month. So: hosts self-schedule on a shared calendar; the chapter thread is member-run (reuse the existing cohort-chat realtime infra, scoped to a `chapter_id`); and ET's only recurring obligation is the monthly "Sparks" review in Studio (15 minutes) plus the occasional new-host intro. No approvals, no event-by-event sign-off. If a host goes quiet, the circle either self-revives or sleeps — and a sleeping circle costs nothing.

Tie to the one metric: every chapter is a retention machine. A member in an active practice circle is the highest-open, highest-click, highest-referral subscriber we have — and each circle is a standing recruiter (members bring craft-friends, who join the Circle to find the circle). Chapters don't dilute the metric; they compound it.

(c) The matching / "your people" mechanic — privacy, consent, safety-first

The single biggest way to wreck trust — and trust is the moat — is to turn "your people" into surveillance. So the mechanic is built on three hard rules, in order of priority: **opt-in, no exposure without consent, safety-first for strangers meeting.**

Rule 1 — Opt-in, granular, reversible. Nothing is shared by default. The profile already has a `visibility` field (private / cohort / public) — we extend it, never weaken it. A member separately opts in to: (a) appearing in near-me counts, (b) the who's-going Atlas markers, (c) being introduced for a practice circle. Each toggle is independent and instantly reversible. Default for all three: **off**.

Rule 2 — Counts before names, double-opt-in before introductions. The product shows aggregates ("3 near you," "11 want to go") — never a named list of strangers a member can browse. There is **no member directory, no map of people, no search-for-humans**. An introduction only happens when *both* sides consent in the same moment: ET says "Maria near you wants to start a ceramics circle — share your first name and contact with her?" and the connection is made *only* if both tap yes. This mirrors widely-used meetup-safety norms — minimise public exposure of who's attending and let people verify before they're identifiable ([Meetup safety practices](#)). Borrow it exactly.

Rule 3 — Safety-first for strangers meeting. Before any first IRL introduction, ET shows a short, plain safety card (not buried in a ToS): meet first in a public place; tell a friend where you're going; you can decline or leave anytime; here's how to report. Hosts get a slightly fuller version. This is non-negotiable and it's cheap — it's a screen of copy, and it signals seriousness, which deepens trust rather than scaring people off.

What the mechanic must NOT be (guardrails, stated plainly so we never drift): not a dating feature; not a "people near you" map; not a follower/friend graph; not a leaderboard of who's most active; not anything that

surfaces a stranger's identity, location, or schedule without that stranger's explicit yes. If a build idea violates any of these, it's out — same as a focus-filter NO, regardless of engagement it might juice.

This restraint is itself the product. In a 2026 where every community app is fighting accusations of surveillance and creep, **“we show you counts, not people, and never introduce anyone without both saying yes”** is a line we can say out loud — and it earns the open, the click, the trust. The honesty lock applies to people, not just Atlas facts.

(d) Alumni & recurring gatherings — the long-term moat

This is the part no incumbent can copy by spending money, because it can only be *grown*. VAWAA sends you home alone; the relationship ends at checkout. MEA's defensible asset isn't its Baja campus — it's the alumni community (a reported 8,000+ across dozens of chapters) that keeps meeting for years, with an annual **Homecoming** as the gravitational centre ([MEA 2026 Homecoming](#)). Those figures are MEA's own and directional; what we adapt is the *shape*, applied honestly to where we actually are.

The honest version of “alumni” (pre-immersion). We have no immersion alumni yet, so we don't pretend to. Our “alumni” today are **circle regulars** — people who've shown up to practice circles repeatedly. We name and honour *that*: a member who's been to several circles is “a regular of the Lisbon ceramics circle,” and that standing is real, earned, and visible on their profile (this is also where the Skill Passport plugs in — see (f)). When the immersion launches in Phase 2, *its* alumni become the deepest tier, but the chapter is the moat whether or not the immersion ever runs.

Recurring gatherings = the rhythm of belonging. Three cadences, light to heavy: - **Monthly:** the practice circle itself (member-run, the heartbeat). - **Seasonal:** an ET-hosted or host-hosted city gathering per cluster — the rung-4 evening, a few times a year, the moment circles in a city meet each other. - **Annual (Phase-2 aspiration, honest):** a single Circle gathering — the equivalent of MEA's Homecoming, adapted to our means. Not a conference, not a festival. One real day at one real source where regulars from several cities convene to make something together. We do not build or promise this now; we name it as the north of the rhythm so chapters know they're building toward something.

Why it's the moat and not just nice: a chapter that has met monthly for two years has a switching cost no competitor can buy — those are *friendships now*. The relationships outlive any single immersion, any algorithm change, any funded rival's ad budget. And every long-running chapter is a self-sustaining source of engaged subscribers and word-of-mouth referrals — the exact compounding the one metric rewards. The moat is people who keep showing up for each other; we just keep introducing them and getting out of the way.

(e) Safety, liability, and trust at solo scale

This is the part that, done wrong, ends ET — a solo founder cannot absorb operator liability for strangers meeting in person. So the legal posture is one sentence, repeated everywhere: **ET introduces; ET does not operate.** ET is the matchmaker and the publisher, never the host, the teacher, the venue, or the tour operator. Everything below enforces that line.

None of the liability statements below are legal advice. They are general guidance drawn from platform help-center pages (Meetup), which are not legal authority. Treat every claim in this section as a starting point to take to a lawyer — see the review TODO at the end of the waiver layer — not as settled law.

The introducer-not-operator stance (load-bearing). The general principle described in meetup-organizer guidance is that an organiser's exposure tends to turn on whether they *organised* the activity — and that risk concentrates on whoever picked the venue, ran the activity, or controlled the event ([Meetup organizer liability](#), [Meetup Liability and Waivers](#)). That is general guidance, not a ruling, and it must be confirmed by a lawyer for our jurisdictions — but it shapes the design: ET deliberately controls *none* of it. The volunteer host picks the place and runs the meet; ET only made an introduction and provided a template. This is why the playbook in (b) is written the way it is — every “ET does not give you: money, insurance, space, or instruction” line is a liability firewall, not just an ops convenience.

The waiver / disclaimer layer (cheap, essential). Platform terms-of-service typically contain a release by members, but — per the same general guidance — such a release may not bind non-member guests a host might bring, so a direct, activity-specific waiver matters whenever there's any physical risk ([Meetup Liability and Waivers](#)). Again: directional, to be confirmed with counsel. Our layer: 1. **A clear disclaimer at opt-in and before every introduction:** “*ET introduces people who share a craft. ET does not organise, supervise, insure, or take responsibility for what members do when they meet. You attend at your own risk and judgement.*” Plain language, not buried. 2. **A host acknowledgement** the volunteer ticks before getting their first intro: they host as an individual, on their own responsibility, are not ET's agent or employee, and ET provides no insurance or indemnity. 3. **An attendee safety + assumption-of-risk card** before a first meet (the (c) card), with explicit assumption of risk for any physical activity (the kind of “higher-risk or unique activity” case where a waiver is generally advised). 4. **Get all of the above reviewed by a lawyer before any rung-3+ goes live.** This is the one place we do not freelance the legal copy, and it is the review that turns every “general guidance” claim above into something we can rely on. Flag it as a real to-do, not a maybe.

What ET must NOT take on (the hard “do not cross” list): - **Do not handle money for IRL meets.** No paid circles, no ticketed gatherings, no “ET takes a cut of the evening.” The moment money flows through ET for an event, ET is arguably an operator. (Phase-2 immersion fees are referral/co-host, arranged **off-site**, and are a different, gated thing.) - **Do not provide or vouch for venues.** Hosts use spaces they already control; ET never rents, inspects, or certifies a location. - **Do not supervise minors, alcohol, or high-hazard activity.** Circles are adults-only by default; any genuinely hazardous craft (open water, fire, blades) carries an extra explicit risk acknowledgement and is the host's domain, not ET's. - **Do not become the safeguarding authority for disputes.** ET sets the safe-space policy and a clear report channel, can remove a member from the Circle, and will — but ET is not the police, the mediator, or the insurer of human behaviour. Be honest about that limit up front.

Trust signals that double as safety: the safe-space policy (adapt the FSF/WordPress template — they exist precisely for this), the report-and-remove path, profile longevity (encourage members to keep a real profile before being introduced — a common meetup vetting heuristic ([Meetup safety](#))), and the counts-not-names default from (c). Each of these protects members *and* signals seriousness — which, again, deepens the engaged-subscriber relationship the one metric measures. Safety done visibly is a trust feature, not a tax.

(f) The loop back — IRL → feed → Skill Passport → more sign-ups

The IRL ladder isn't a cost centre we tolerate; it's the **richest content and proof engine we have**, and it closes a self-reinforcing loop straight back into the one metric. Three returns from every real-world moment:

1. IRL → feed content. The single thing ET asks every host and attendee for is one photo or one 30–60s clip per meet, tagged **#FromTheCircle**. This is gold for the Daily Drop: it's wordless, it's hands-on-real-material, it ends on a finished object or a moment of competence — it is the locked format, made by real members instead of staged. It's also the most honest possible answer to the master test “does it fight the screen or feed it?” — the footage exists *because* people put the screen down and met. We don't pivot the Daily Drop format (stop-doing list); we *feed* it from circles. A steady stream of #FromTheCircle clips means the feed sustains itself without Arnaud filming everything at sea.

2. IRL → Skill Passport evidence. The community already has a **prior-experiences** table with an instructor-VERIFIED flag, and the future build target is verifiable skill credentials. The practice circle is where self-declared becomes *witnessed*. When a member attends a circle, the **host can confirm attendance** (a light, honest tap — “yes, they were here and they threw a pot”), and that confirmation becomes a real entry on the member's Skill Passport: “*3 sessions · Lisbon ceramics circle · host-confirmed.*” Not a fabricated badge — a witnessed fact, which honours the honesty lock (verified against a real human who was there, exactly like Atlas facts are verified against the source). This is the credential layer made trustworthy *by* the IRL layer. It's also the bridge to Phase-2: when the immersion gate opens, a member's circle history is the proof they're ready — “earned, not bought,” recorded.

3. Proof → more Circle sign-ups. #FromTheCircle clips in the feed and on Atlas pages show prospects the thing the funded incumbents can't show: *real peers, really gathered, really making something*. That's the highest-converting recruitment we have — it's social proof of “your people” that no listicle competitor can fake. Each clip's caption carries the one Atlas page URL and the one-line “what we are,” driving viewers to the Circle to find their own circle. And every new sign-up feeds rung 1, which seeds the next chapter — closing the loop.

The full circuit: a near-me signal → a practice circle → a witnessed Skill Passport entry + a #FromTheCircle clip → that clip recruits the next member → who raises their hand near-me → who seeds the next circle. IRL doesn't drain the one metric. It is the flywheel that makes the one metric self-sustaining.

What we build now vs later

Now-moves (no money, no operations, builds on what's live): - Rung 1 (near-me count) and Rung 2 (who's-going Atlas markers) — extend the existing profile **interests/skills/visibility** fields; counts-not-names; opt-in off by default. - The host playbook page + message/safe-space templates on **community.html**. - The “Sparks” view in Studio (city × craft cluster ranking) — a 15-min/month founder review. - Chapter threads — reuse cohort-chat infra scoped to **chapter_id**. - #FromTheCircle clip ask → Daily Drop pipeline; host-confirmed attendance → Skill Passport entries. - The consent + safety + disclaimer copy (and the lawyer review before rung-3 goes live).

Phase-2-gated later-moves (only when all 8 gates are green): - Rung 5, the cohort immersion — off-site, referral/co-host fee, never re-surfaced as a public checkout. - The annual Homecoming-style gathering — named as the north of the rhythm, not promised.

The line we never cross: no money through ET for IRL meets; no people-directory or location map; no introduction without both-sides consent; no ET-as-operator. Cross any of those and we trade the moat (honesty, trust, focus) for engagement we don't want. Hold them, and the IRL ladder becomes the single strongest producer of engaged Circle subscribers we have — because the people on it aren't subscribers anymore. They're friends who happen to get our email.

Sources (third-party, directional — verify in Studio Metrics before relying on any figure): [MEA Wisdom Community](#) · [MEA alumni portal](#) · [MEA Homecoming 2026](#) · [WordPress Meetup Program Basics](#) · [FSF LibreLocal 2026](#) · [Meetup organizer liability](#) · [Meetup Liability and Waivers](#) · [Meetup safety practices](#). Liability links are help-center pages, not legal authority — general guidance only, not legal advice.

Part 7 — Development vectors: the different ways to grow

Everything below feeds one number: **engaged Circle subscribers** — sign-ups that actually open and click (~150–300 quality subs by day 90, 40%+ open). Not revenue, not raw followers, not pageviews. A vector that doesn't move that number, or can't be traced to it within a step or two, is not a now-move. And every vector passes the same gate it asks of immersions: **embodied, immersive, community-based, real** — and the master test, *does it fight the screen or feed it?*

The Atlas is finished. We have 99 disciplines, 324 destinations, 835 verified school cards, 424 SEO pages. **Distribution is the work now, not page count.** Read this whole section as a menu of how to distribute, deepen, and own — never as a reason to build more rooms in a house nobody's walked into yet.

***A note on the numbers in this section.** Any external engagement figure you see anywhere in this guideline is from third-party 2026 analyses — directional, not gospel. Treat them as hypotheses to confirm in the Studio Metrics tab against our own Plausible and Resend data, never as settled law. The honesty posture is the moat; it has to hold internally too.*

How to read the tags

Tag	Meaning
NOW	Build/run in the next 90 days. Directly grows engaged Circle subs.
NEXT	Months 2–6, once the NOW vectors are running and measured.
PHASE-2-GATED	Touch nothing here until all 8 gates are true. Listing it is planning, not permission.

The 8 gates, restated so they're never forgotten: (1) several hundred non-branded organic visitors/mo for 3+ months; (2) ~1,000–2,000 engaged subs; (3) 40%+ open / 6–9% click sustained, with replies and referrals; (4) regular unprompted “can you connect me?” inbound; (5) a few clips genuinely travelled; (6) ≥1 vetted partner willing to co-host on referral; (7) a soft refundable-deposit off-site test fills; (8) founder capacity — one immersion won't starve the Atlas and the Circle.

The vectors, in priority order

1. Content & SEO depth — the Atlas + Journal compounding · NOW

Thesis. The Atlas is the only asset that earns trust while Arnaud is at sea. 424 pages already exist; their job is to be *found*, then to hand the visitor a reason to join the Circle. The Journal (science-and-stories essays — San Sebastián is live) is the “why step up” pillar that turns a searcher into a believer. SEO is the cheapest, most durable feeder of non-branded organic visitors — gate (1) and gate (4) live here.

Smallest honest first step. Submit `sitemap.xml` in Google Search Console (Arnaud’s standing TODO), then add **one Journal essay every two weeks** from the 8 briefs already seeded in Studio (Lambert/effort, Csikszentmihalyi/flow, the dive-reflex, Ericsson/Newport, SDT, Waldinger/Blue Zones, motor-learning). Each essay links to its matching Atlas page and ends on one Circle CTA.

Leading indicator. Non-branded organic sessions/week in Plausible (already live site-wide), and Circle sign-ups attributed to `/atlas/` and `/journal/` landing pages. Watch GSC impressions → clicks on the discipline/destination pages.

Trap to avoid. Building *more* Atlas. 99 disciplines is the ceiling per the STOP-DOING list. Also: writing essays that *feed the screen* (engagement-bait listicles) — every Journal piece must point off-screen, toward a place and a person. Never the “embodied mastery vs the AI age” billboard; it’s a quiet one-line foil, not a headline (buzzword fatigue).

Gate. None — pure NOW. This is the foundation the 8-gate is measured against.

2. The social engine — the Daily Drop · NOW

Thesis. One locked format — 30–60s wordless vertical process clip, hands and raw material, ending on the finished object — is our reach engine. It fights the screen by *showing* what a screen can’t teach. Same clip, same day, to TikTok + Reels + Shorts on @educatedtraveler.app. The caption carries the only conversion that matters: one Atlas page URL and the one-line “what we are.”

Smallest honest first step. Ship the first 10 from the Studio Ideas tab (`SHOT-LIST-first10.md`) on **founder-shot or licensed stock footage** — decoupled from cold school-licensing so it survives charter legs. Cadence floor: **3×/week + a 3–4 week scheduled buffer**. Get the IG Business account + Meta token so `NEXUS/Tools/social-publish` can auto-post.

Leading indicator. Saves and shares (not raw views) per clip; click-through from caption to the Atlas page; and the only one that counts — Circle sign-ups that arrived via a clip. Saves > likes is the signal that someone is *planning to go*. (Per 2026 third-party analyses, shares carry outsized distribution weight relative to likes — treat that as directional and confirm it in Studio Metrics, not as a platform constant.)

Trap to avoid. Pivoting the format before day 90 (locked). Chasing view-count vanity. Adding words/voiceover/captions-on-screen — wordless is the differentiator. Posting clips that don’t end on competence.

Gate. None — NOW. Clips genuinely travelling is gate (5), so this engine also unlocks Phase 2.

3. The Circle as owned audience · NOW — this is the metric itself

Thesis. The Circle is the one audience we own, not rent from an algorithm. It is the north-star number incarnate. The sender is deployed (Resend, three issues written, welcome auto-fires on `launch_waitlist` insert). The whole job now is: more quality sign-ups, higher open/click, more replies.

Smallest honest first step. Wire the `launch_waitlist` → `circle-welcome` database webhook (the one remaining activation step), send the first real broadcast of issue-01, and make the homepage `#circle` form the single most prominent CTA on every Atlas and Journal page (it already is the CTA — verify it converts). Ask one question in every issue that invites a reply (“which craft is calling you?”).

Leading indicator. Open rate (target 40%+), click rate (6–9%), and **reply count** — replies are the truest proxy for the unprompted “can you connect me?” inbound that gate (4) demands. Track unsubscribe rate as the honesty check.

Trap to avoid. Buying or padding the list — a 60% open of 200 real people beats a 5% open of 5,000. Broadcasting before the welcome flow is wired (you’d burn first impressions). Treating the Circle as a sales channel — it is a campfire, never a checkout.

Gate. None — NOW. Gates (2) and (3) ARE the Circle’s own metrics scaled up.

4. The instructor / school network · NOW (relationship) → PHASE-2 (money)

Thesis. The 835 verified school cards are 835 warm-able relationships. Schools want the footage, the link, the audience. Recruiting them as *content partners* now (footage, a quote, a fact-check) is the same motion that, much later, recruits them as *immersion co-hosts* — gate (6). The OUTREACH.md email already doubles as Phase-2 partner recruitment.

Smallest honest first step. Use the existing outreach scripts (`next-outreach-batch.py`, `outreach-digest.py`, `find-followups.py`) to send the footage-licensing / fact-confirm email to the **highest community-rank schools per core**, ~4 at a time (the concurrency lesson from Wave 2). Ask only for footage and a fact-check — no money, no co-host pitch yet.

Leading indicator. Reply rate from schools; footage clips received; and any unprompted “can our students join your Circle?” — the early shape of gate (6).

Trap to avoid. Re-surfacing the instructor-availability portal or any booking plumbing (STOP-DOING). Pitching co-hosting before the audience exists — a school says no to an empty room. Presenting any unverified claim as fact in outreach (honesty lock — it’s the moat with exactly these people). Never let any co-host or fee language leak into the rung-1 footage email; the money rung stays gated and out of public-facing copy.

Gate. Relationship-building is NOW. Any **referral/co-host fee is PHASE-2-GATED** (all 8, especially gate 6 — a vetted partner willing to co-host).

5. The community feed — skills-only social surface · NEXT

Thesis. A feed where people chasing the same craft share learning clips, facts, and progress turns the Circle from a broadcast list into a place people *return to* — raising engagement (gate 3) and seeding the “your people” promise. The profile/auth/cohort-chat machinery already exists; this is an additive surface, not a new platform.

Smallest honest first step. Don't build the feed yet. **First prove demand inside the Circle:** in a broadcast, invite subscribers to reply with what they're learning, then republish the best (with permission) in the next issue — a hand-run “feed” with zero code. If replies are strong, build the lightest possible public feed on the existing Supabase profiles.

Leading indicator. Reply-with-content rate to the “what are you learning?” prompts; if built, weekly posters and weekly returners (not registrations).

Trap to avoid. Building the feed before there's an audience to fill it — an empty feed is worse than none. Letting it slide into the skillcation-listicle swamp or a generic social timeline. It must stay *skills-only* and pass the focus filter, or it feeds the screen.

Gate. NEXT — only after the Circle is replying. Not gated by the 8 (it's free), but gated by demonstrated demand.

6. The Skill Passport — verifiable credentials • NEXT (light) → PHASE-2 (issuance)

Thesis. Credential is one-third of the defensible fusion (Credential + Curation + Cohort) no incumbent stacks. VAWAA gives you no certificate; we eventually could. Today “credentials” are self-declared free text; only `prior-experiences` carries an instructor-VERIFIED flag — that verified flag is the honest seed of a Skill Passport.

Smallest honest first step. Do **nothing new in code.** Lean on the existing instructor-verified `prior-experiences` flag. The first “passport” is simply: a profile that shows *verified* prior experiences distinctly from self-declared ones. Watch whether members ask to have an experience verified — that's the demand signal.

Leading indicator. Number of members requesting verification of a prior experience; instructors willing to verify.

Trap to avoid. Issuing badges/certificates before they mean anything — an unearned badge is the gamification we retired forever (XP/level/badge bans). A credential that isn't hand-verifiable betrays the honesty lock. Never let it become a vanity sticker. Any credential-standard fact this depends on (Open Badges / Verifiable Credentials status, LinkedIn's Add-to-Profile policy) must be re-verified at build time, not trusted from a dated note — these change.

Gate. Light surfacing of the existing verified flag = NEXT. Real **issuance is PHASE-2-GATED** — it only has meaning once immersions run and instructors can vouch for what someone actually did at the source.

7. IRL circles — the who's-going / local-meetup mechanism • NEXT

Thesis. “Your people” has to leave the screen eventually or it's a lie. A lightweight who's-going / local-circle signal is the bridge from owned audience to real gathering — and it's the strongest possible proof of intent ahead of any immersion.

Smallest honest first step. Manual, not built: in a Circle issue, ask “where are you, and what are you chasing?” and **hand-match two people** in the same city chasing the same craft. One real introduction is worth more than a meetup feature.

Leading indicator. Number of accepted member-to-member introductions; any self-organised meetup that happens because of a Circle prompt.

Trap to avoid. Building a who's-going/events feature on speculation. Liability and safety of organising IRL gatherings — keep it introduction-only, off-platform, until volume justifies structure (and get any organiser-liability posture reviewed by a lawyer before a single hosted gathering — comparison-site help pages are general guidance, not legal authority). Don't let it drift into generic "travel buddy" matching (fails the focus filter).

Gate. NEXT — manual matching now, any feature only after the pattern is real. Genuine inbound intros relate to gate (4).

8. Partnerships & cross-promotion · NOW (Littoralicious) → NEXT (complementary brands)

Thesis. Arnaud already runs Littoralicious — a live publication with an audience and the same warmth (never the same CSS/fonts). The culinary core of the Atlas overlaps Littoralicious's world directly. A clearly-labelled, voice-respecting crossover is free, warm reach to exactly the right people. Complementary brands (dive shops, craft schools, sail lofts already in the Atlas) extend this later.

Smallest honest first step. One Littoralicious crossover: a Littoralicious piece on a craft-at-the-source theme (e.g. New Basque / San Sebastián, which the Journal already covers) with a single honest line pointing to the EducatedTraveler Atlas page and Circle. No banner-swap, no co-branding mess — one warm referral, both directions, kept editorially separate.

Leading indicator. Circle sign-ups attributed to the Littoralicious referral; click-through on the crossover link.

Trap to avoid. Bleeding the two brands' CSS/fonts/identity together (design lock — same warmth, never shared stylesheet). Cross-promoting anything that fails the focus filter just because it's adjacent. Treating Littoralicious's list as a list to *sell* to — it's a list to *introduce* to.

Gate. Littoralicious crossover = NOW (owned, free, on-brand). Outside complementary-brand partnerships = NEXT, and any of them carrying money = PHASE-2-GATED.

9. Honest monetisation · PHASE-2-GATED

Thesis. Money comes *later* — Phase 2, after all 8 gates — and only in ways that never put a price or a checkout on the public surface. The public surface stays sell-nothing; that's the moat. There is **one** sanctioned revenue shape, and a single far-off idea that stays parked behind the same gate.

- **Curated partner immersions (the sanctioned one) — PHASE-2-GATED.** Pick 2–3 trusted schools from the Atlas, brand them as EducatedTraveler experiences, take a **referral / co-host fee, arranged OFF-SITE, never an on-site transaction.** Launch as a single product page + soft refundable deposit (the *barcelona-vermut* page is the intentional prototype shape), NOT the dormant marketplace flow. **All 8 gates true first — no exceptions.**
- **Supporter membership (parked, not a now/next move) — PHASE-2-GATED, out of scope until the gate.** The canon confines money to Phase-2 curated partner immersions after the 8-gate. So treat any "support the Atlas" patronage idea as **deferred behind the same 8-gate, not a possible**

NEXT — do not stand up a donation or support surface on the public site before the gate is met, and never let it read as a paywall or a gate on any content. If it ever happens at all, it happens *after* the gate, and only if every Atlas fact stays free and the label is unambiguous (never “unlock”). Until then it is parked: build nothing, promise nothing.

Smallest honest first step. Nothing on either. For immersions: wait for the gates. For supporter membership: it is parked behind the gate too — the only thing to *notice* in the meantime is whether engaged members unprompted *ask* how to support the work; even that is just a signal to log, not a trigger to build.

Leading indicator. For immersions: gate (4) inbound “can you connect me?” and gate (6) a willing co-host. (Supporter membership has no now-metric — it is gated.)

Trap to avoid. Re-surfacing any of the dormant selling machinery — prices (\$3.5k–\$40k), Quest selector, XP/confetti/badges, PayPal/Stripe checkout, enrollment flows — on the public site (STOP-DOING + public-surface bans). Putting a price anywhere public. Standing up any support/donation surface before the 8-gate. Inventing a TAM to justify any of it (the \$75B–\$432B figures were refuted — never cite them).

Gate. Immersions = **all 8 gates**. Supporter membership = **also all 8 gates, parked until then** — and only ever if it never compromises the sell-nothing public surface.

If you only do five things in the next 90 days

Five moves, each chosen because it directly grows **engaged Circle subscribers**. Do these; let everything else wait.

1. **Wire the Circle welcome webhook and send the first real broadcast.** `launch_waitlist` → `circle-welcome`, then ship issue-01. This is the metric itself — there is no excuse for the sender to sit deployed-but-silent. (*Vector 3.*)
2. **Submit the sitemap in GSC and publish one Journal essay every two weeks.** Turn the 424 finished pages and the 8 seeded essays into found, believed, off-screen pull — the cheapest durable feeder of non-branded visitors and Circle sign-ups. (*Vector 1.*)
3. **Run the Daily Drop at 3×/week from founder/stock footage, with the IG Business account + Meta token live.** Ship the first 10 from the Ideas tab. Every caption ends on one Atlas URL; every Atlas page ends on the Circle. (*Vector 2.*)
4. **Warm 4 high-rank schools per core for footage and fact-checks** (≤ 4 at a time). Same motion that becomes Phase-2 co-host recruitment — builds gate (6) relationships while we have zero pricing to defend. (*Vector 4.*)
5. **Run one Littoralicious crossover into the Atlas + Circle**, and inside every Circle issue ask one reply-inviting question (“which craft is calling you?” / “where are you, what are you chasing?”). Free warm reach in; replies out — feeding gates (2), (3), and (4) at once. (*Vectors 8, 3, 7.*)

Each of the five ends at the same campfire: a person who opens, clicks, and replies. That is the whole job for 90 days. **A place, a person, your people. Earned, not bought.**

Part 8 — The Tram Rail: intelligent, honest automation

A tram runs fast because it’s locked to a rail. It can’t wander into the crowd. That’s the whole idea here. Automation should carry the dull weight — file-shuffling, scheduling, first drafts, reminders — so Arnaud’s scarce hours go to the two things no machine can fake: **verifying a fact** and **answering a human warmly**. Every minute the rail saves has exactly one job: more engaged Circle subscribers. If a piece of automation doesn’t feed that metric, or it touches a fact or a person’s trust, it doesn’t get full speed. It gets a human at the controls.

The discipline that makes this safe is already in the codebase. `ig-lint.py` has a `BANNED` list and a credential-keyword gate. `build-daily-drop.py` refuses banned words and prices at generation time. `circle-welcome` / `circle-broadcast` / `circle-unsubscribe` already send mail through Resend. We are not building an empire. We are wiring the gaps and putting a linter in front of the door.

***A note on the figures and policy claims below.** The platform rules, governance patterns, and “2026 best-practice” splits cited here are drawn from third-party 2026 analyses and provider help-pages — they are directional, not gospel, and they shift. Verify the load-bearing ones (platform disclosure rules, model id/pricing) against the primary source at build time, and track real performance in the Studio Metrics tab rather than betting the format on someone’s blog number.*

(a) The Automation Map — one level per surface

Three levels, and the rule for choosing is blunt: **anything factual or trust-bearing can never be full-auto**. Atlas facts, who-teaches-what, credentials — a machine may *draft* the claim, but a human asserts it. That’s the HONESTY LOCK rendered as plumbing.

Surface	Level	Why this level
Atlas fact-keeping (school cards, masters, courses, confidence tags)	Human-only	This is the moat. Every fact is hand-verified against the school’s official source. A model may pre-fill a “provisional, verifying” row, but Arnaud confirms before the confidence tag flips. Never auto-assert who teaches what.
/atlas page builds (the 424 static SEO pages)	Full-auto	<code>build-atlas-pages.py</code> renders verified data into HTML. The facts were human-gated upstream; the rendering is deterministic and safe to run on every deploy. Mechanical, not editorial.
Daily-Drop scheduling (calendar, which atlas page each clip points at)	Full-auto	<code>build-daily-drop.py</code> already selects by communityRank, round-robins the 4 cores, no repeats, and refuses banned words. The selection is rules-based, so let it run.
Daily-Drop clip itself (the 30-60s wordless process video)	Human-only	The clip is the soul. Real hands, real material, shot or sourced by Arnaud. Never AI-generated footage — that fails the master test (“fight the screen or feed

Surface	Level	Why this level
		it?”) and would be a lie about “at the source.”
Caption / hook / idea generation	Draft-then-approve	A model drafts the caption from the locked recipe (superlative hook + atlas ‘why’ + community/season + one-line “what we are” + the one URL + 4-5 hashtags). <code>ig-lint.py</code> gates it. Arnaud reads it before it posts.
Journal drafting (the essay pillar)	Draft-then-approve	A model can assemble a cited outline and a rough draft from verified sources. But the Journal is “science-and-stories” with citations — every claim needs a human check, and the voice must be founder-led, so Arnaud rewrites and signs it.
Instructor outreach (first contact)	Draft-then-approve	<code>next-outreach-batch.py</code> picks the queue; a model drafts the personalised note. A real human relationship starts here — Arnaud reads and edits each one. Never blast. Connect, don’t sell.
Outreach follow-ups	Draft-then-approve	<code>find-followups.py</code> + <code>outreach-digest.py</code> already surface who’s overdue. The <i>reminder</i> is auto; the <i>message</i> is human-approved. A follow-up that sounds like a bot ends the relationship.
Circle welcome email	Full-auto	<code>circle-welcome</code> fires on sign-up with pre-written, locked copy. It’s transactional and the text is fixed. Safe to run untouched.
Circle broadcast (the issues)	Draft-then-approve	<code>circle-broadcast</code> is the rail; the <i>issue content</i> is written/edited by Arnaud and lint-checked before send. This is the one-metric surface — open rate and replies live or die on it. Never auto-publish an issue.
Community-feed ops & moderation (the future skills-only feed)	Draft-then-approve for action, human-only for judgment	A model can flag (spam, banned words, off-topic, abuse) and queue. A human decides removals and bans. Never auto-delete a member’s post on a model’s say-so; never auto-generate “engagement.”
Credential issuance (future verifiable badges)	Human-only	The hardest lock. A credential is a promise that someone can do a real thing. It is issued only after a human (Arnaud or a vetted instructor) verifies — exactly like the existing instructor-VERIFIED flag on prior-experiences. A model never mints a badge.

Surface	Level	Why this level
Metrics (Plausible + Circle open/click + outreach digest)	Full-auto	Reading and summarising numbers is safe and high-leverage. Auto-pull into the Studio Metrics tab. The <i>interpretation</i> and any decision stays human.

The pattern, in one line: **render, schedule, remind, summarise** → **full-auto**. **Write, message, judge** → **draft-then-approve**. **Assert a fact, mint a credential, generate “people”** → **human-only**. This mirrors what 2026 third-party analyses describe as mature ops — automating the mechanical majority of the work and keeping humans on tone, framing, and the publish gate (directional, per [ContentMonk pattern via Parseur](#); Zapier on HITL patterns).

(b) Build on what exists — wire next vs. leave manual

Wire next (cheap, high-leverage, all build on live pieces):

1. **Extend `ig-lint.py` into a universal pre-publish gate** (see §c). It already has the **BANNED** list and credential check — point it at Daily-Drop captions, Journal drafts, *and* Circle broadcasts, not just IG drafts. One linter, every door.
2. **Caption drafter wired to `build-daily-drop.py`**. The calendar already names the discipline, destination, communityRank, and URL per day. Feed that row to a model with the locked caption recipe; write the draft into the Studio **Daily Drop** tab as “needs approval.” Saves ~3-5 min/clip × daily.
3. **Outreach drafter wired to `next-outreach-batch.py` + `find-followups.py`**. The queue and the overdue list exist; add a draft note per row into the Studio **Outreach** tab. Saves ~10-15 min/contact.
4. **Metrics auto-pull into the Studio Metrics tab**. Plausible API + Circle open/click + `outreach-digest.py` → one weekly snapshot. Saves the manual dashboard-checking, ~20 min/week, and keeps THE ONE METRIC visible.
5. **Finish NEXUS/Tools/social-publish (`ig_publish.py`)** once the IG Business account + Meta token land. That closes the loop from approved caption → posted, across TikTok/Reels/Shorts. It’s already built and waiting on credentials, not code.

Leave manual (on purpose):

- **The Atlas itself**. STOP-DOING is explicit: 99 disciplines / 424 pages is plenty. No automated discovery of new schools. Distribution is progress, not page count.
- **The clip and the Journal voice**. Founder hands, founder pen.
- **Any reply to a real person** in the Circle or the feed. The whole product is that a human answers. Automating replies kills the thing being sold.
- **Anything from the dormant selling machinery** (PayPal/Stripe edge fns, enrollment, prices). Those functions exist in the repo; they stay dormant. Never wire them to the public site.

This keeps the solo-founder stack lean. The pattern 2026 third-party analyses point to — content-creation + scheduling + summarisation as the automatable layers, with the founder on judgment — is exactly the split above (directional, per [Reinventing.ai on solo operators](#)).

(c) Honesty guardrails — what makes the automation safe

These are non-negotiable. They are the HONESTY LOCK as enforced rules, not good intentions.

1. **Human-in-the-loop on everything factual or trust-bearing.** Atlas facts, credentials, and who-teaches-what are *never* auto-asserted. A model may draft; a human asserts. If a fact can't be verified against the school's official source, it ships as **"provisional, verifying"** with a low-confidence tag, never as fact. This is the "provisional, verifying" discipline written into the workflow, not just the copy.
2. **No fake engagement, ever.** No bots, no sockpuppets, no purchased followers, no AI-generated "members" in the Circle or the feed, no auto-liking. Community-RANK is "best chance to find your cohort," not reputation theatre. Faked engagement would corrupt THE ONE METRIC (engaged subscribers) at its root — you'd be optimising a lie.
3. **AI-assistance disclosure norms.** The clip is real footage, always — so there's nothing synthetic to disclose there, which is the honest position and, per current 2026 platform guidance, the compliant one (the read at the time of writing is that TikTok/Meta require disclosure only for *synthetic or significantly AI-altered realistic* media, with real process footage exempt — but treat this as a moving rule and re-verify against the platforms' own policy pages before relying on it: [platform policies summary](#); [TikTok synthetic-media rules](#)). Where a model *did* draft copy that a human edited, that reads as standard assisted writing under current norms — but the rule we hold is simpler and stricter than any platform's: **we never present a machine-made claim as a verified human one.** No AI-generated faces, voices, or fabricated "students." If we ever publish anything genuinely synthetic, we label it plainly. Trust is the product ([FTC disclosure guidance](#); [creator authenticity checklist](#)).
4. **The brand-lock / banned-word linter as a hard gate before anything posts.** Extend `ig-lint.py` so *nothing reaches a public surface without passing it*. Concretely:
 - Keep the existing BANNED list; add the canon's public-surface bans — `book`, `enroll`, `checkout`, `Quest`, `XP/level/badge` (as gamification), `price/$`, and "choose your depth" as a pricing ladder.
 - Add a positive check: the verb must be **introduce / connect / bridge**, never `sell / book / enroll`.
 - Add a **fact-claim flag**: if a draft asserts a specific master, school, or credential, the linter requires a `verified: true` field in frontmatter (sourced) or it fails. A claim with no source can't pass.
 - Wire it as a pre-commit / pre-publish step in `build-daily-drop.py`, the Journal build, and the `circle-broadcast` send path. Exit 1 = no post. The linter is the tram's rail-guard: copy cannot derail into hype, selling, or an unsourced fact.

The principle 2026 governance write-ups converge on: automation is safe only with **audit trails and a human publish gate** (directional, per [cflow on 2026 trends](#)). The linter is our audit trail — every pass/fail is logged.

(d) The Model-Refresh Loop — keep the rail modern without breaking the locks

Models improve fast; the rail should ride the best available without ever loosening a lock. **Default to the latest, most capable model when building any AI feature** — for ET's drafting (captions, Journal outlines, outreach notes) that means the current frontier Claude model rather than an older or cheaper tier, because the cost is trivial against a solo founder's hours and quality directly feeds open/reply rates. Verify the exact current model id and pricing against Anthropic's docs at build time — never hardcode a model id or price from memory into a script or this doc.

The ritual — "Tram Service" — runs monthly (last Sunday, ~30 min), recorded in the Studio Plan tab under a MODEL-REFRESH note and mirrored to the MEMORY project file:

Checklist: - [] **What's new?** Check for a newer/more capable model from the providers we use (Claude first). Note the current best model id + pricing from the provider's own docs. - [] **Re-bench the drafters.** Run the same three real prompts (one caption, one outreach note, one Journal paragraph) through the current model

vs. the new one. Does the new one need fewer edits? Fewer linter fails? - [] **Cost check.** Is the new model affordable at our volume (a few dozen drafts/week)? Almost always yes — quality wins. - [] **Locks still hold?** After switching, run the full `ig-lint.py` suite on the new outputs. If the new model trips the linter more, that's a regression — fix the prompt or hold. - [] **One change at a time.** Swap the model behind *one* surface (e.g. captions) first, watch one week of linter pass-rate and edit-time, then roll to the others. Never swap everything blind. - [] **New tool scan.** Any new automation tool worth wiring (better scheduler, better transcription for clip captions)? Only adopt if it saves real founder-minutes and respects the offline-tolerance rule. - [] **Record it.** Date, model id, what changed, edit-time delta, linter pass-rate. One line in the Studio Plan note so the next refresh has a baseline.

The locks (design, voice, banned words, honesty, no-selling) are **upstream of the model**. The model is a swappable engine; the rail-guard (`ig-lint.py`) doesn't care which engine drafted the text. That's what lets you adopt a better model in minutes without risking the brand.

(e) Pragmatic build order for a solo founder at sea (offline-tolerant)

Everything here is **draft-offline, send-when-connected** — the founder is often at sea, so generation happens locally (scripts + cached data + the model), and only the *send* steps need a connection. Time saved is per week, the unit that matters.

1. **Extend `ig-lint.py` into the universal gate** (½ day build). *Saves: nothing directly — it's the seatbelt that makes everything below safe to automate.* Do this first; it's the rail.
2. **Caption drafter on `build-daily-drop.py` → Studio Daily-Drop tab** (½ day). Generates a week of lint-passed caption drafts in one offline run; Arnaud approves in batches. *Saves ~30-40 min/week.*
3. **Metrics auto-pull → Studio Metrics tab** (½ day). One weekly snapshot of THE ONE METRIC (engaged subs, open/click) + outreach digest. *Saves ~20 min/week and keeps focus on the right number.*
4. **Outreach drafter on `next-outreach-batch.py` + `find-followups.py`** (1 day). Pre-drafted, personalised, lint-passed notes queued for approval. *Saves ~60-90 min/week once outreach is regular.*
5. **Journal draft-assist** (½ day prompt work, no new infra). Cited outline + rough draft from verified sources into the Articles tab. *Saves ~2-3 hrs per essay — but the rewrite stays human.*
6. **Finish `social-publish` / `ig_publish.py`** (gated on IG Business + Meta token, not on code). Closes approved-caption → posted across all three platforms. *Saves ~15-20 min/day of manual cross-posting.*

Phase-gated (do NOT build now): the skills-only feed's moderation automation, credential issuance, and the linktree/IRL mechanisms are later-moves — they only matter once the Circle is large enough to need them (toward the PHASE-2 GATE). Building moderation before there's a feed, or credential-minting before there are credentials, is gold-plating. The rail carries today's load: get the clip out, the caption right, the issue sent, the outreach warm, the number visible. That's what grows the one metric.

The whole rail in a sentence: **let machines render, schedule, remind, and draft — let Arnaud verify and connect — and put one linter in front of every door.** Fast where it's safe, human where it counts.

Sources (third-party 2026 analyses — directional, verify before betting the format on them): Parseur — Future of HITL AI 2026 · Zapier — Human-in-the-loop patterns · cflow — AI workflow trends 2026 · Reinventing.ai — solo operators 2026 · susit.ai — platform AI policies 2026 · auditsocials — TikTok synthetic-media rules · humanadsai — FTC AI disclosure · crescitaly — AI creator authenticity checklist

Part 9 — Sequencing & the focus filter

The nine Parts are a menu, not a to-do list. If we chase all of them at once we lose focus, and losing focus is the one failure this whole document exists to prevent. So here is the order — what we do NOW, what comes NEXT, and what stays locked behind the 8-gate until the numbers earn it. Everything below is judged by one question: does it move **engaged Circle subscribers**? If a move can't be traced to that number within a step or two, it waits.

NOW (the next 90 days — the only fight)

These run in parallel because they are one loop: attention in, sign-ups out.

- **The Daily Drop, every day, format locked** (Part 2). The wordless process clip earns the eyeball. It is the top of the funnel; nothing else works without it. Do not pivot the format pre-day-90.
- **Every public surface routes to one Atlas URL or the Circle** (Parts 1, 2). One destination per piece. Caption carries the “what we are” line and a single link. *Maps to the metric*: every view that clicks through is a sign-up candidate.
- **The relationship program, one discipline at a time** (Part 3). Feature schools, send them warm strangers, earn the repost. We introduce; the vouch is never for sale. *Maps to the metric*: every school share puts our line and our link in front of exactly the people most likely to open and click.
- **Studio Metrics is the scoreboard** (Parts 2, 7, 8). Read our own per-stream data weekly. Treat every external figure as a hypothesis to confirm here, never as a law. *Maps to the metric*: we steer the format toward what actually produces engaged subs.

NEXT (start small while NOW is running — deepen retention, don't replace the funnel)

These make the people we already have open more, reply more, refer more. They layer on; they never pull effort off the Daily Drop.

- **The Skill Passport — profile first** (Part 5). Turn the read-only crew card at /crew into one public page a person is proud to put in their bio. Build the profile now; build the credential model second and slower. The passport's only job is to make the honesty visible. *Maps to the metric*: a member proud of their page shares it — and a shared page is a warm new sign-up.
- **The Community room — skills-only, signal over dopamine** (Part 4). A room where people chasing the same fire show each other their hands at work, not another feed. Small and active beats big and empty. *Maps to the metric*: a member who replies and refers is the deepest kind of engaged sub.
- **The Tram Rail — wire the gaps** (Part 8). Let automation carry the dull weight so Arnaud's scarce hours go to the two things no machine can fake: verifying a fact and answering a human warmly. Keep the linter in front of the door. *Maps to the metric*: every hour the rail saves is an hour spent producing more engaged subs.

PHASE-2-GATED (do not start until all 8 are true)

- **IRL — the first off-site connection** (Part 6). A click becomes a coffee, a coffee a practice circle. The soft refundable-deposit immersion test (gate 7) lives here. IRL is not a new product; it is the deepest expression of the one metric.

- **Curated partner immersions — the first money** (Parts 3, 7). Referral or co-host fee, arranged off-site, never an on-site transaction, only after the 8-gate. One immersion at a time, and only if it won't starve Atlas + Circle.

The month-6 checkpoint (kill / persevere)

At day 180, look at the real numbers in Studio Metrics — not the third-party figures, ours. Persevere if the trend toward the 8-gate is real: organic non-branded traffic climbing toward “several hundred a month, 3+ months sustained,” Circle subs growing with 40%+ open and 6–9% click, and the first unprompted “can you connect me?” notes landing in the inbox. If the metric is flat after six honest months of the locked format and one-discipline-at-a-time outreach, do not add features and do not pivot the video — change the *angle and the hook*, the cheapest variable, and run another clean stretch. We kill effort that doesn't move the one number; we do not kill the line.

Before you ship anything

Run the Focus Filter, then remember the **STOP-DOING list** above — most of the ways we can lose are things we already decided not to do. And remember the first law, the law under every other line in this document: **never lose focus**. Open on the payoff, one idea per piece, re-hook, close on one thing. Every line earns the next, or it goes.

Appendix — One-page quick reference

Print this. It is the whole Guideline at a glance.

The spine (say these). “A place, a person, your people.” · “Earned, not bought.” · “At the source, not the simulation.” · “Skills last, tans fade.” · The Atlas (the map) · The Circle (the community). Customer is the hero; we are the guide. Fight the screen; honour the source; people first.

THE ONE METRIC. Engaged Circle subscribers — sign-ups that open and click. ~150–300 quality subs by day 90, 40%+ open. Not revenue, not followers, not pageviews.

Voice — do. Warm, plain, founder-led. Open on the payoff. One idea per section. Re-hook the middle. Close on one thing. The verb is INTRODUCE / CONNECT / BRIDGE.

Voice — don't. No hype. Never sell / book / enroll. Never cite an unverified market-size or TAM number. Never present an unverified fact as fact — mark it “provisional, verifying,” low-confidence. No decorative emoji.

Banned words. transformation · transformational · life-changing · vacation · holiday · luxury · easy. Public-surface bans: book · enroll · checkout · prices · “Quest” · XP/levels/badges · “choose your depth” as a pricing ladder.

Focus Filter (any NO = out). Embodied? · Immersive (a week+)? · Community-based? · Returns a real skill? · Does it fight the screen or feed it?

Design tokens. Warm-Dark Editorial. bg ink #0d0b09 · surface #14110d · text paper #f3ede2 · muted rgba(243,237,226,.56) · hairline rgba(243,237,226,.09). Accents: SEA #7fa8a5 · EMBER #d28a52. Cores: wellness #94ad86 · adventure #6fa3a0 · creative #cf8f6e · culinary #c9a24a. Fonts: Fraunces · Inter · IBM Plex Mono. Same warmth as Littoralicious, never its CSS or fonts.

The five things to do in the next 90 days. 1. Ship the Daily Drop every day, format locked — wordless process clip ending on the finished object, one place-name card, ambient sound, same clip to TikTok + Reels + Shorts. 2. Route every public surface to one destination — a single Atlas URL or the Circle — carrying the “what we are” line. 3. Run the relationship program one discipline at a time — feature schools, send warm strangers, earn the repost; introduce, never sell. 4. Build the Skill Passport profile (not credentials yet) from the existing /crew card — a page people are proud to share. 5. Read Studio Metrics weekly and steer to what actually produces engaged subs — confirm every external figure against our own data first.

Model-refresh ritual. This document is reviewed on a fixed cadence so the line never drifts: - **Weekly** — read Studio Metrics; check the one number; retire any angle that isn’t producing engaged subs. - **Monthly** — re-read The Line and the STOP-DOING list before planning the next batch; confirm nothing has crept back onto the public site. - **Day 90 / Day 180** — the kill-or-persevere checkpoint against the trend toward the 8-gate (Part 9). - **Update the Last reviewed date** at the top whenever any LOCK changes, and never let a Part contradict The Line — if it does, The Line wins.

First law, under everything: **never lose focus.** Every line earns the next.